


CSE 258 – Lecture 3

Web Mining and Recommender Systems

Supervised learning – Classification

Regularization

Regularization is the process of penalizing model complexity during training


$$\arg \min_{\theta} = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$


How much should we trade-off accuracy versus complexity?

Optimizing the (regularized) model

$$\arg \min_{\theta} = \underbrace{\frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2}_{f(\theta)}$$

- Could look for a closed form solution as we did before
- Or, we can try to solve using **gradient descent**



Optimizing the (regularized) model

Gradient descent:

1. Initialize θ at random
2. While (not converged) do
$$\theta := \theta - \alpha f'(\theta)$$

All sorts of annoying issues:

- How to initialize theta?
- How to determine when the process has converged?
- How to set the step size alpha

These aren't really the point of this class though

Optimizing the (regularized) model

$$f(\theta) = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$

$$\frac{\partial f}{\partial \theta_k} ?$$

$$f(\theta) = \frac{1}{N} \sum_i (y_i - \underbrace{X_{i \cdot} \theta}_{\sum_k X_{ik} \theta_k})^2 + \lambda \sum_k \theta_k^2$$

$$\frac{\partial f}{\partial \theta_k} = \frac{1}{N} \sum_i -2 X_{ik} (y_i - X_{i \cdot} \theta) + 2 \lambda \theta_k$$


Optimizing the (regularized) model

Gradient descent in scipy:

(code for all examples is on <http://jmcauley.ucsd.edu/cse258/code/week1.py>)

(see “ridge regression” in the “sklearn” module)

Model selection

$$\arg \min_{\theta} = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$


How much should we trade-off accuracy versus complexity?

Each value of lambda generates a different model. **Q:** How do we select which one is the best?

Model selection

How to select which model is best?

A1: The one with the lowest training error?

A2: The one with the lowest test error?

We need a **third** sample of the data that is not used for training or testing

Model selection

A **validation set** is constructed to "tune" the model's parameters


- Training set: used to **optimize the model's parameters** → choose 0
- Test set: used to report how well we expect the model to perform on **unseen data** → only used once
- Validation set: used to **tune** any model parameters that are not directly optimized → choose 1

Model selection

A few “theorems” about training, validation, and test sets

- The training error **increases** as lambda **increases**
- The validation and test error are at least as large as the training error (assuming infinitely large random partitions)
- The validation/test error will usually have a “sweet spot” between under- and over-fitting

Model selection



Summary of Week 1: Regression

- Linear regression and least-squares
 - (a little bit of) feature design
 - Overfitting and regularization
 - Gradient descent
- Training, validation, and testing
 - Model selection

Homework

Homework is **available** on the course webpage

<http://cseweb.ucsd.edu/classes/fa18/cse258-a/files/homework1.pdf>

Please submit it by the beginning of the **week 3** lecture (Oct 15)

All submissions should be made as **pdf files on gradescope**

CSE 258 – Lecture 3

Web Mining and Recommender Systems

Recap


Last week...

Last week we started looking at
supervised learning problems

$$f(\text{data}) \xrightarrow{?} \text{labels}$$

Last week...

We studied **linear regression**, in order to learn linear relationships between features and parameters to predict **real-valued** outputs



Last week...



ratings
features

Product Details

Genres	Science Fiction, Action, Horror
Director	David Twohy
Starring	Vin Diesel, Radha Mitchell
Supporting actors	Cole Hauser, Keith David, Lewis Fitz-Gerald, Claudia Black, Rhiana Granger, Angela Moore, Peter Chiang, Ken Twohy
Studio	NBC Universal
MPAA rating	R (Restricted)
Captions and subtitles	English Details ▼
Rental rights	24 hour viewing period. Details ▼
Purchase rights	Stream instantly and download to 2 locations Details ▼
Format	Amazon Instant Video (streaming online video and digital download)

$f(\text{user features}, \text{movie features}) \xrightarrow{?} \text{star rating}$

Four important ideas from last week:

1) Regression can be cast in terms of **maximizing a likelihood**

$$y_i = X_i \cdot \theta + \mathcal{N}(0, \sigma)$$
$$\max_{\theta} \prod_i \underbrace{p_{\theta}(y_i | X_i)} \leadsto \min_{\theta} \sum_i (y_i - X_i \cdot \theta)^2$$

Four important ideas from last week:

2) Gradient descent for model optimization


1. Initialize θ at random
2. While (not converged) do

$$\theta := \theta - \alpha f'(\theta)$$

Four important ideas from last week:

3) Regularization & Occam's razor

Regularization is the process of penalizing model complexity during training

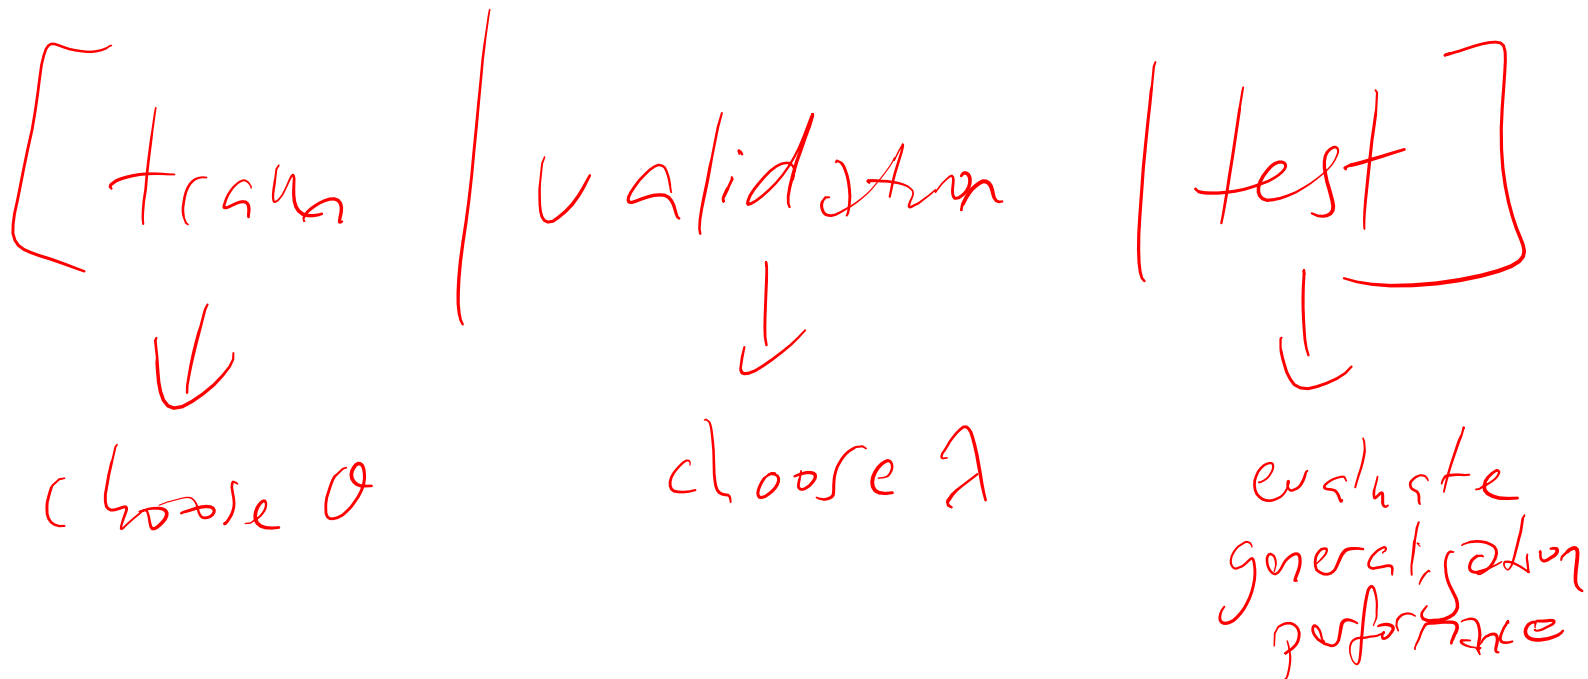
$$\arg \min_{\theta} = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$


How much should we trade-off accuracy versus complexity?

Four important ideas from last week:

4) Regularization pipeline

1. Training set – select model parameters
2. Validation set – to choose amongst models (i.e., hyperparameters)
3. Test set – just for testing!



Model selection

A **validation set** is constructed to “tune” the model’s parameters

- Training set: used to **optimize the model’s parameters**
- Test set: used to report how well we expect the model to perform on **unseen data**
- Validation set: used to **tune** any model parameters that are not directly optimized


Model selection

A few “theorems” about training, validation, and test sets

- The training error **increases** as lambda **increases**
- The validation and test error are at least as large as the training error (assuming infinitely large random partitions)
- The validation/test error will usually have a “sweet spot” between under- and over-fitting

Today...

How can we predict **binary** or **categorical** variables?











Today...



Will I **purchase**
this product?
(yes)

Shop for engagement rings on Google Sponsored ⓘ

 French-Set Halo Diamond Eng... \$1,990.00 Ritani	 18K White Gold Delicate... \$950.00 Brilliant Earth ★★★★★ (57)	 18K White Gold Fancy D... \$1,825.00 Brilliant Earth ★★★★★ (13)	 Chamise Diamond Eng... \$975.00 Brilliant Earth ★★★★★ (7)
 Vintage Cushion Halo... \$4,140.00	 Princess Cut Diamond Eng... \$1,906.82	 18K White Gold Hudson... \$975.00	 18K White Gold Harmon... \$1,675.00

Will I **click on**
this ad?
(no)

Today...

What animal appears in this image?
(mandarin duck)



Today...

What are the **categories** of the item being described?

(book, fiction, philosophical fiction)

From [Booklist](#)

Houellebecq's deeply philosophical novel is about an alienated young man searching for happiness in the computer age. Bored with the world and too weary to try to adapt to the foibles of friends and coworkers, he retreats into himself, descending into depression while attempting to analyze the passions of the people around him. Houellebecq uses his nameless narrator as a vehicle for extended exploration into the meanings and manifestations of love and desire in human interactions. Ironically, as the narrator attempts to define love in increasingly abstract terms, he becomes less and less capable of experiencing that which he is so desperate to understand. Intelligent and well written, the short novel is a thought-provoking inspection of a generation's confusion about all things sexual. Houellebecq captures precisely the cynical disillusionment of disaffected youth. *Bonnie Johnston* --This text refers to an out of print or unavailable edition of this title.

Today...

We'll attempt to build **classifiers** that make decisions according to rules of the form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

This week...

1. Naïve Bayes

Assumes an **independence** relationship between the features and the class label and “learns” a simple model by counting

2. Logistic regression

Adapts the **regression** approaches we saw last week to binary problems

3. Support Vector Machines

Learns to classify items by finding a hyperplane that separates them

This week...

Ranking results in order of how likely they are to be relevant

A screenshot of a Google search interface. At the top, a search bar contains the text "tea station". Below the search bar, navigation tabs are visible: "Web" (highlighted with a red underline), "Maps", "Shopping", "Images", "News", "More" (with a dropdown arrow), and "Search tools". Below the tabs, the search results summary shows "About 20,900,000 results (0.61 seconds)". The first search result is for "Tea Station 加州茶棧" with the URL "teastationusa.com/". The snippet describes 12 locations in California and Nevada and expresses gratitude to tea lovers. It includes a 3.8-star rating from 19 Google reviews and links to "Write a review" and "Google+ page". A location pin icon is followed by the address "7315 Clairemont Mesa Boulevard, San Diego, CA 92111" and the phone number "(858) 268-8198". Below the address are links for "Menu", "About", "Ten Ren Products", and "San Gabriel". At the bottom, there is a link for "Tea Station - Kearny Mesa - San Diego, CA | Yelp" and a snippet from "www.yelp.com" showing "Restaurants > Chinese" and a "Yelp" link.

tea station

Web Maps Shopping Images News More Search tools

About 20,900,000 results (0.61 seconds)

Tea Station 加州茶棧
teastationusa.com/ ▼
12 Tea Station locations in California and Nevada making Tea Station the ... We'd like to take this moment to thank you all tea lovers for your continued support.
3.8 ★★★★★ 19 Google reviews · Write a review · Google+ page



7315 Clairemont Mesa Boulevard, San Diego, CA 92111
(858) 268-8198
Menu - About - Ten Ren Products - San Gabriel

Tea Station - Kearny Mesa - San Diego, CA | Yelp
www.yelp.com > Restaurants > Chinese > Yelp

This week...

Evaluating classifiers

- False positives are nuisances but false negatives are disastrous (or vice versa)
 - Some classes are very rare
- When we only care about the “most confident” predictions



e.g. which of these bags contains a weapon?

Naïve Bayes

We want to associate a probability with a label and its negation:




$$\begin{array}{c} p(label|data) \\ p(\neg label|data) \\ = | \end{array}$$

(classify according to whichever probability is greater than 0.5)

Q: How far can we get just by counting?

Naïve Bayes

e.g. $p(\text{movie is "action"} \mid \text{schwarzenegger in cast})$



Just count!

#fims with Arnold = 45

#**action** films with Arnold = 32

$p(\text{movie is "action"} \mid \text{schwarzenegger in cast}) = 32/45$

$$p(a|y) = \frac{p(a,y)}{p(y)}$$

Naïve Bayes

What about:

$p(\text{movie is "action" |}$
 schwarzenneger in cast **and**
 release year = 2017 **and**
 mpaa rating = PG **and**
 budget < \$1000000
)

#(training) fims with Arnold, released in 2017, rated PG, with a
 budged below \$1M = 0

#(training) action fims with Arnold, released in 2017, rated PG,
 with a budged below \$1M = 0

Naïve Bayes

Q: If we've never seen this combination of features before, what can we conclude about their probability?

A: We need some **simplifying assumption** in order to associate a probability with this feature combination

Naïve Bayes

Naïve Bayes assumes that features are **conditionally independent** given the label

$$(feature_i \perp\!\!\!\perp feature_j | label)$$

Naïve Bayes

$(feature_i \perp\!\!\!\perp feature_j | label)$

$$\times p(a, b) = p(a)p(b) \text{ (independence)}$$

$$p(a, b | c) = p(a | c)p(b | c)$$

$a =$ I'm wearing shorts

$b =$ you're wearing shorts

$c =$ It's 100° (or 40°C)

Conditional independence?

$$(a \perp\!\!\!\perp b | c)$$

(a is conditionally independent of b, given c)

“if you know **c**, then knowing
a provides no additional
information about **b**”

(I remembered my umbrella $\perp\!\!\!\perp$ the streets are wet | it's raining)

Naïve Bayes

$$(feature_i \perp\!\!\!\perp feature_j | label)$$



$$p(feature_i, feature_j | label)$$

=

$$p(feature_i | label)p(feature_j | label)$$

Naïve Bayes

posterior



$$p(\text{label}|\text{features})$$

$$p(A|B) = \frac{p(A)p(B|A)}{p(B)}$$

"Bayes Rule"

cond
indep

prior



likelihood



$$= p(\text{label}) \underbrace{p(\text{features}|\text{label})}_{p(\text{features})}$$

evidence

$$p(\text{label}) \frac{\prod_i p(\text{feature}_i|\text{label})}{p(\text{features})}$$

Naïve Bayes

$$p(\text{label}|\text{features}) = \frac{p(\text{label}) \prod_i p(\text{feature}_i|\text{label})}{\underbrace{p(\text{features})}_{?}}$$

$$\frac{p(\text{label}|\text{features})}{p(\neg\text{label}|\text{features})}$$

The denominator doesn't matter, because we really just care about

$$p(\text{label}|\text{features}) \quad \text{vs.} \quad p(\neg\text{label}|\text{features})$$

both of which have the same denominator

Naïve Bayes

The denominator doesn't matter, because we really just care about

$$p(\textit{label}|\textit{features}) \quad \textbf{vs.} \quad p(\neg \textit{label}|\textit{features})$$

both of which have the same denominator

Example 1

Amazon editorial descriptions:

Amazon.com Review

For most children, summer vacation is something to look forward to. But not for our 13-year-old uncle, and cousin who detest him. The third book in J.K. Rowling's [Harry Potter series](#) catapults Dursleys' dreadful visitor Aunt Marge to inflate like a monstrous balloon and drift up to the ceiling (and from officials at Hogwarts School of Witchcraft and Wizardry who strictly forbid students to go out into the darkness with his heavy trunk and his owl Hedwig).

As it turns out, Harry isn't punished at all for his errant wizardry. Instead he is mysteriously rescued by a triple-decker, violently purple bus to spend the remaining weeks of summer in a friendly inn called the Leaky Cauldron. This book explains why the officials let him off easily. It seems that Sirius Black is loose. Not only that, but he's after Harry Potter. But why? And why do the Dementors, the guard dogs of the Ministry of Magic, are unaffected? Once again, Rowling has created a mystery that will have children and adults clamoring for the next book. Fortunately, there are four more in the works. (Ages 9 and older) --Karin Snelson --This text refers to the paperback edition.

50k descriptions:

http://jmcauley.ucsd.edu/cse258/data/amazon/book_descriptions_50000.json

Example 1

P(book is a children's book |
 "wizard" is mentioned in the description **and**
 "witch" is mentioned in the description)

Code available on:

<http://jmcauley.ucsd.edu/cse258/code/week2.py>

Example 1

Conditional independence assumption:

“if you know **a book is for children**, then knowing that **wizards are mentioned** provides no additional information about whether **witches are mentioned**”

obviously ridiculous

Double-counting

Q: What would happen if we trained two regressors, and attempted to “naively” combine their parameters?

Double-counting

$$\begin{aligned} \text{Weight} &= 20 \text{ height} \\ &= 15 (\text{shoe size}) \end{aligned}$$

$$\text{Weight} = 20 \text{ height} + 15 (\text{shoe size})$$

Double-counting

A: Since both features encode essentially the same information, we'll end up **double-counting** their effect

Logistic regression

Logistic Regression also aims
to model

$$p(\textit{label}|\textit{data})$$

By training a classifier of the
form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Logistic regression

Last week: regression

$$y_i = X_i \cdot \theta$$

This week: logistic regression


$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Logistic regression

Q: How to convert a real-valued expression ($X_i \cdot \theta \in \mathbb{R}$)
Into a probability
($p_\theta(y_i|X_i) \in [0, 1]$)

Logistic regression

A: sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$



Logistic regression

Training:

$X_i \cdot \theta$ should be maximized
when y_i is positive and
minimized when y_i is
negative

$$\begin{aligned} \arg \max_{\theta} & \prod_{y_i=1} p_{\theta}(y_i | X) \prod_{y_i=0} (1 - p_{\theta}(y_i | X)) \\ &= \prod_{y_i=1} \sigma(X_i \cdot \theta) \prod_{y_i=0} (1 - \sigma(X_i \cdot \theta)) \end{aligned}$$

Logistic regression

How to optimize?

$$L_{\theta}(y|X) = \prod_{y_i=1} p_{\theta}(y_i|X_i) \prod_{y_i=0} (1 - p_{\theta}(y_i|X_i))$$

- Take logarithm
- **Subtract** regularizer
- Compute gradient
- Solve using gradient **ascent**

Logistic regression

$$L_{\theta}(y|X) = \prod_{y_i=1} p_{\theta}(y_i|X_i) \prod_{y_i=0} (1 - p_{\theta}(y_i|X_i))$$

$$\begin{aligned} \ell_{\theta}(y|X) &= \sum_{y_i=1} \log \sigma(X_i \cdot \theta) + \sum_{y_i=0} \log (1 - \sigma(X_i \cdot \theta)) \\ &= \sum_{y_i=1} \log \left(\frac{1}{1 + e^{-X_i \cdot \theta}} \right) + \sum_{y_i=0} \log \left(\frac{e^{-X_i \cdot \theta}}{1 + e^{-X_i \cdot \theta}} \right) \\ &= \sum_{\textcircled{y_i}} -\log(1 + e^{-X_i \cdot \theta}) + \sum_{y_i=0} -X_i \cdot \theta \end{aligned}$$

Logistic regression

$$l_{\theta}(y|X) = \sum_i -\log(1 + e^{-\underbrace{X_i \cdot \theta}}) + \sum_{y_i=0} -X_i \cdot \theta - \lambda \|\theta\|_2^2$$

$$\frac{\partial l}{\partial \theta_k} =$$

$$- \sum_i x_{ik} \theta_k$$

$$\sum_i \frac{x_{ik} e^{-X_i \cdot \theta}}{1 + e^{-X_i \cdot \theta}} + \sum_{y_i=0} -x_{ik} - 2\lambda \theta_k$$

$$= \sum_i x_{ik} (1 - \sigma(X_i \cdot \theta)) + \sum_{y_i=0} -x_{ik} - 2\lambda \theta_k$$

Logistic regression

Log-likelihood:

$$l_{\theta}(y|X) = \sum_i -\log(1 + e^{-X_i \cdot \theta}) + \sum_{y_i=0} -X_i \cdot \theta - \lambda \|\theta\|_2^2$$

Derivative:

$$\frac{\partial l}{\partial \theta_k} = \sum_i X_{ik} (1 - \sigma(X_i \cdot \theta)) + \sum_{y_i=0} -X_{ik} - 2\lambda \theta_k$$

Multiclass classification

The most common way to generalize **binary** classification (output in $\{0,1\}$) to **multiclass** classification (output in $\{1 \dots N\}$) is simply to train a binary predictor for each class

e.g. based on the description of this book:

- Is it a Children's book? {yes, no}
- Is it a Romance? {yes, no}
- Is it Science Fiction? {yes, no}
- ...

In the event that predictions are inconsistent, choose the one with the highest confidence

Questions?

Further reading:

- On Discriminative vs. Generative classifiers: A comparison of logistic regression and naïve Bayes (Ng & Jordan '01)
- Boyd-Fletcher-Goldfarb-Shanno algorithm (BFGS)


CSE 258 – Lecture 3

Web Mining and Recommender Systems

Supervised learning – SVMs


Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?



Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?



Logistic regression

- Logistic regressors don't optimize the number of "mistakes"
- No special attention is paid to the "difficult" instances – every instance influences the model
- But "easy" instances can affect the model (and in a bad way!)
- How can we develop a classifier that optimizes the number of mislabeled examples?

Support Vector Machines

This is essentially the intuition behind Support Vector Machines (SVMs) – train a classifier that focuses on the “difficult” examples by minimizing the misclassification error

We still want a classifier of the form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta - \alpha > 0 \\ -1 & \text{otherwise} \end{cases}$$

But we want to minimize the number of misclassifications:


$$\arg \min_{\theta} \sum_i \delta(y_i(X_i \cdot \theta - \alpha) \leq 0)$$

Support Vector Machines


$$\arg \min_{\theta} \sum_i \delta(y_i(X_i \cdot \theta - \alpha) \leq 0)$$

Support Vector Machines

Simple (seperable) case: there exists a perfect classifier




Support Vector Machines




The classifier is defined by the hyperplane $\theta \mathbf{x} - \alpha = 0$

Support Vector Machines



Q: Is one of these classifiers preferable over the others?

Support Vector Machines




A: Choose the classifier that maximizes the distance to the nearest point

Support Vector Machines

Distance from a point to a line?

Support Vector Machines



$$\arg \min_{\theta, \alpha} \frac{1}{2} \|\theta\|_2^2$$

such that

$$\forall_i y_i (\theta \cdot X_i - \alpha) \geq 1$$

Support Vector Machines

This is known as a
"quadratic program" (QP)
and can be solved using
"standard" techniques

$$\arg \min_{\theta, \alpha} \frac{1}{2} \|\theta\|_2^2$$


such that

$$\forall_i y_i (\theta \cdot X_i - \alpha) \geq 1$$

See e.g. Nocedal & Wright ("Numerical Optimization"), 2006


Support Vector Machines

But: is finding such a separating hyperplane even possible?




Support Vector Machines

Or: is it actually a good idea?



Support Vector Machines

Want the margin to be as wide as possible



While penalizing points on the wrong side of it

Support Vector Machines


Soft-margin formulation:

$$\arg \min_{\theta, \alpha} \quad \frac{1}{2} \|\theta\|_2^2$$

such that

$$\forall_i y_i (\theta \cdot X_i - \alpha) \geq 1$$

Judging a book by its cover



[0.723845, 0.153926, 0.757238, 0.983643, ...]

4096-dimensional image features

Images features are available for each book on
http://jmcauley.ucsd.edu/cse258/data/amazon/book_images_5000.json



<http://caffe.berkeleyvision.org/>

Judging a book by its cover

Example: train an SVM to predict whether a book is a children's book from its cover art

(code available on)

<http://jmcauley.ucsd.edu/cse258/code/week2.py>

Judging a book by its cover

- The number of errors we made was extremely low, yet our classifier doesn't seem to be very good – why?
(stay tuned next lecture!)

Summary

The classifiers we've seen today all attempt to make decisions by associating weights (θ) with features (x) and classifying according to

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Summary

- **Naïve Bayes**

- Probabilistic model (fits $p(\text{label}|\text{data})$)
- Makes a conditional independence assumption of the form $(\text{feature}_i \perp\!\!\!\perp \text{feature}_j | \text{label})$ allowing us to define the model by computing $p(\text{feature}_i | \text{label})$ for each feature
- Simple to compute just by counting

- **Logistic Regression**

- Fixes the “double counting” problem present in naïve Bayes

- **SVMs**

- Non-probabilistic: optimizes the classification error rather than the likelihood

Questions?