CSE 258 – Lecture 4

Web Mining and Recommender Systems

Evaluating Classifiers

Last lecture...

How can we predict **binary** or **categorical** variables?

$$\begin{tikzcd} & \{0,1\}, \{\text{True, False}\} \\ f(\text{data}) \overset{?}{\to} \text{labels} \arrow[ur] \arrow[dr] & \\ & \{1, \dots, N\} \end{tikzcd}$$

Last lecture...

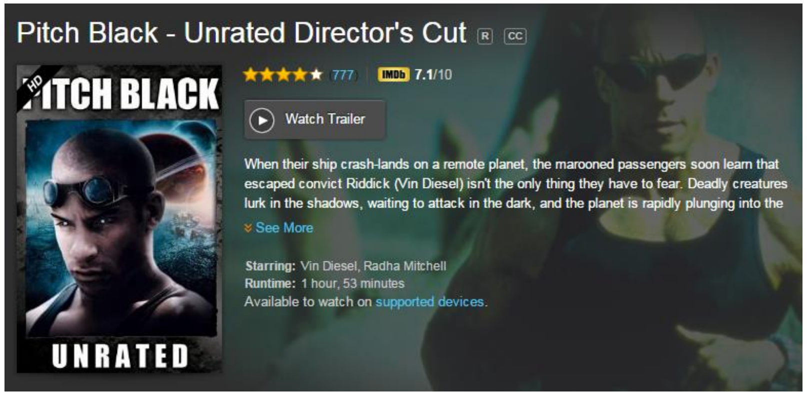

Pitch Black - Unrated Director's Cut

★★★★★ 777 IMDb 7.1/10

Watch Trailer

When their ship crash-lands on a remote planet, the marooned passengers soon learn that escaped convict Riddick (Vin Diesel) isn't the only thing they have to fear. Deadly creatures lurk in the shadows, waiting to attack in the dark, and the planet is rapidly plunging into the

See More

Starring: Vin Diesel, Radha Mitchell
Runtime: 1 hour, 53 minutes
Available to watch on supported devices.

Will I **purchase** this product?

(yes)

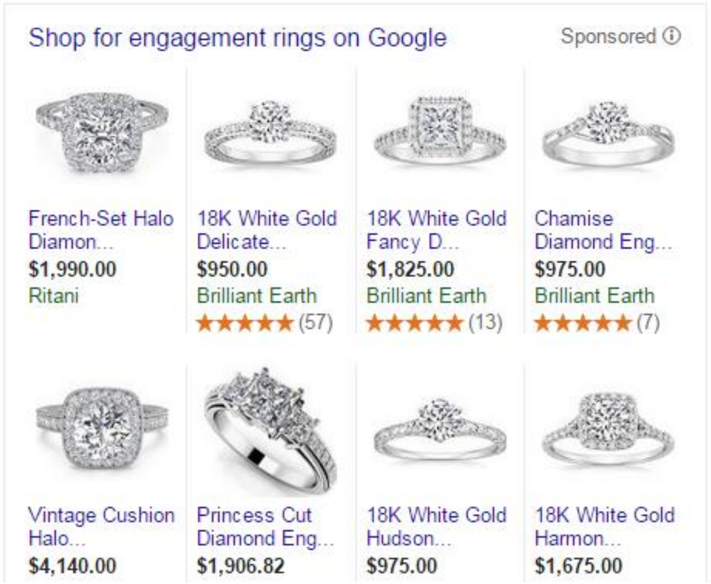

Shop for engagement rings on Google — Sponsored

French-Set Halo Diamon... $1,990.00 Ritani	18K White Gold Delicate... $950.00 Brilliant Earth ★★★★★ (57)	18K White Gold Fancy D... $1,825.00 Brilliant Earth ★★★★★ (13)	Chamise Diamond Eng... $975.00 Brilliant Earth ★★★★★ (7)
Vintage Cushion Halo... $4,140.00	Princess Cut Diamond Eng... $1,906.82	18K White Gold Hudson... $975.00	18K White Gold Harmon... $1,675.00

Will I **click on** this ad?

(no)

Last lecture...

- **Naïve Bayes**
  - Probabilistic model (fits $p(label|data)$)
  - Makes a conditional independence assumption of the form $(feature_i \perp\!\!\!\perp feature_j | label)$ allowing us to define the model by computing $p(feature_i|label)$ for each feature
  - Simple to compute just by counting
- **Logistic Regression**
  - Fixes the "double counting" problem present in naïve Bayes
- **SVMs**
  - Non-probabilistic: optimizes the classification error rather than the likelihood

1) Naïve Bayes

$$\overset{\text{posterior}}{\overset{\downarrow}{p(label|features)}} = \frac{\overset{\text{prior}}{\overset{\downarrow}{p(label)}}\,\overset{\text{likelihood}}{\overset{\downarrow}{p(features|label)}}}{\underset{\text{evidence}}{\underset{\uparrow}{p(features)}}}$$

due to our conditional independence assumption:

$$p(label|features) = \frac{p(label)\prod_i p(feature_i|label)}{p(features)}$$

2) logistic regression

sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$

$\frac{1}{1+e^{-X_i \cdot \theta}}$

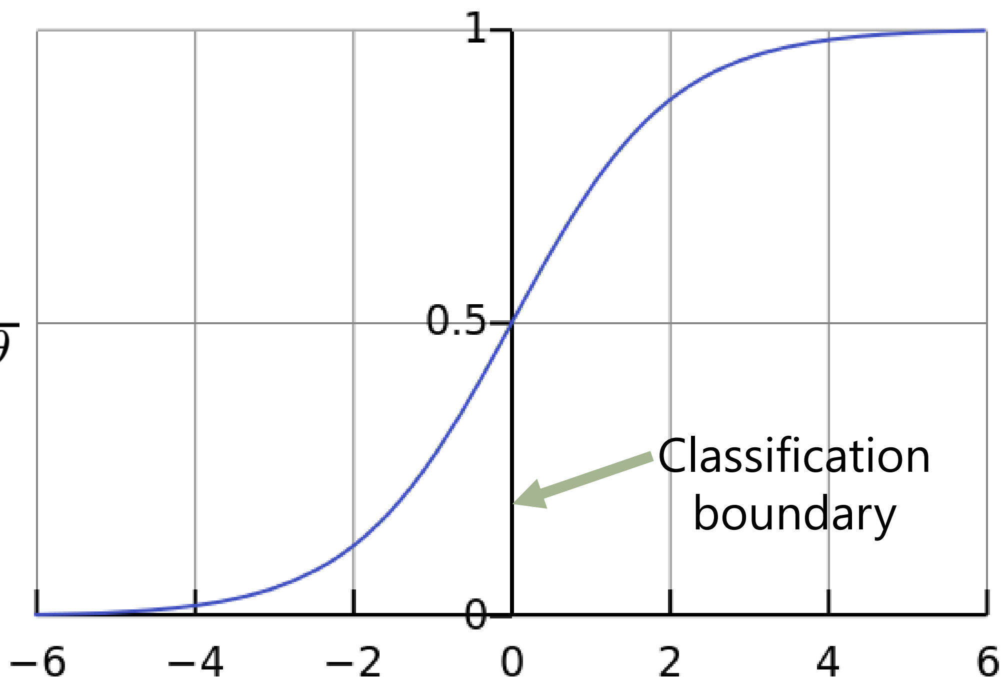

$X_i \cdot \theta$

Logistic regression

- Logistic regressors don’t optimize the number of “mistakes”
- No special attention is paid to the “difficult” instances – every instance influences the model
- But “easy” instances can affect the model (and in a bad way!)
- How can we develop a classifier that optimizes the number of mislabeled examples?

Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?

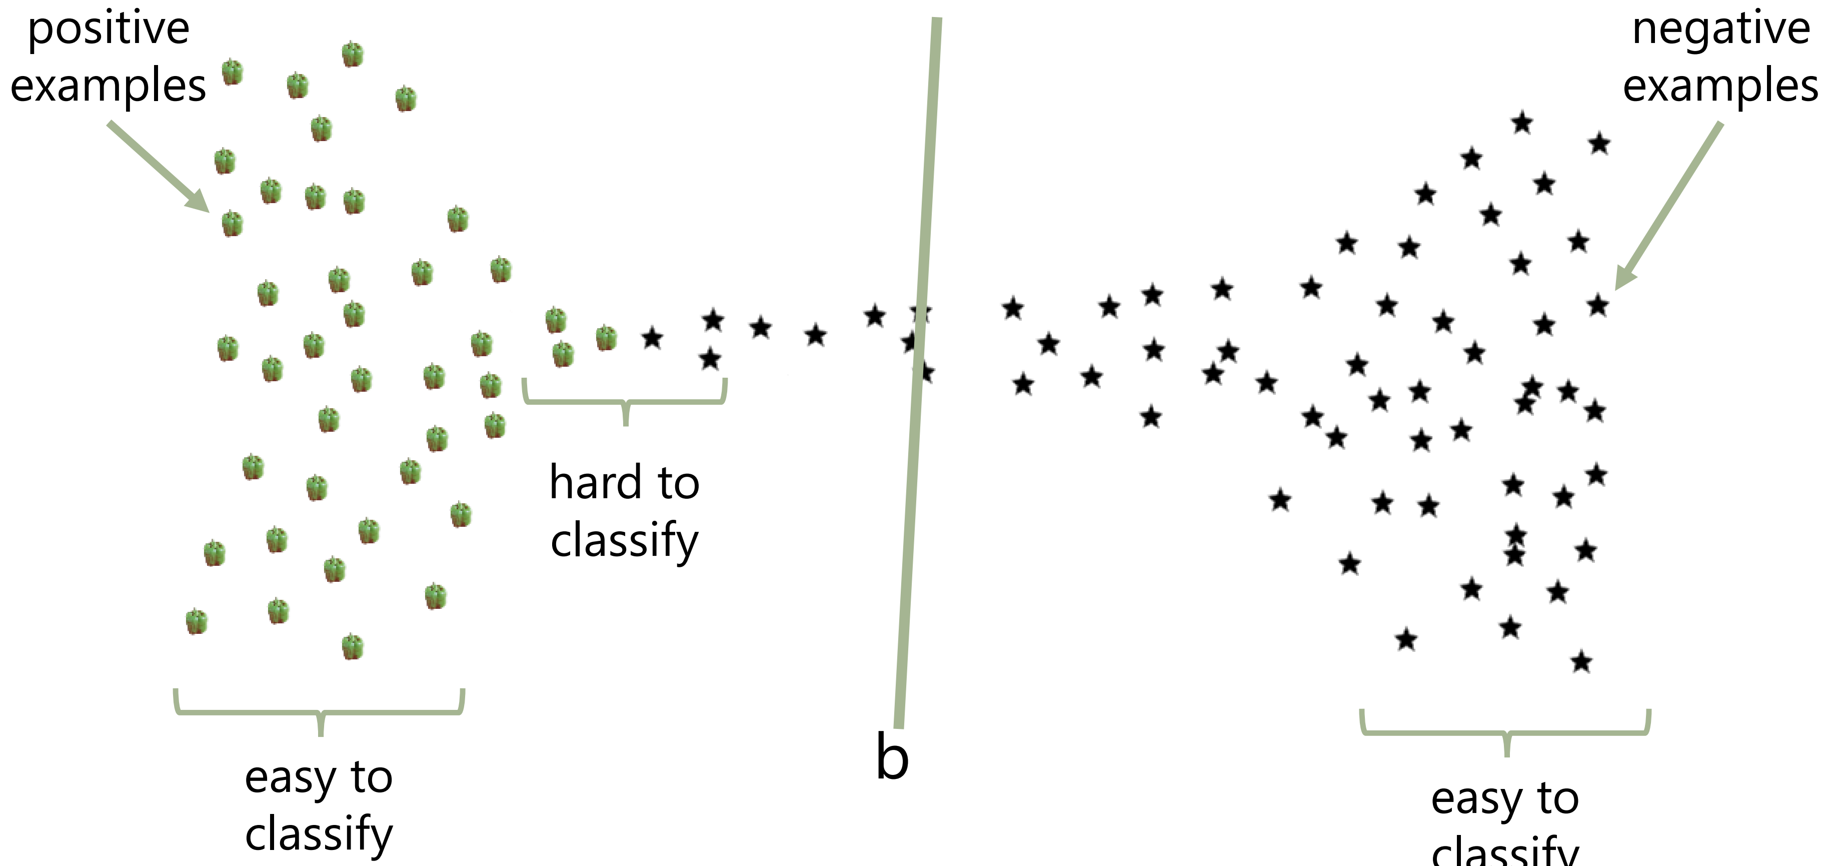

3) Support Vector Machines

Try to optimize the **misclassification error** rather than maximize a probability

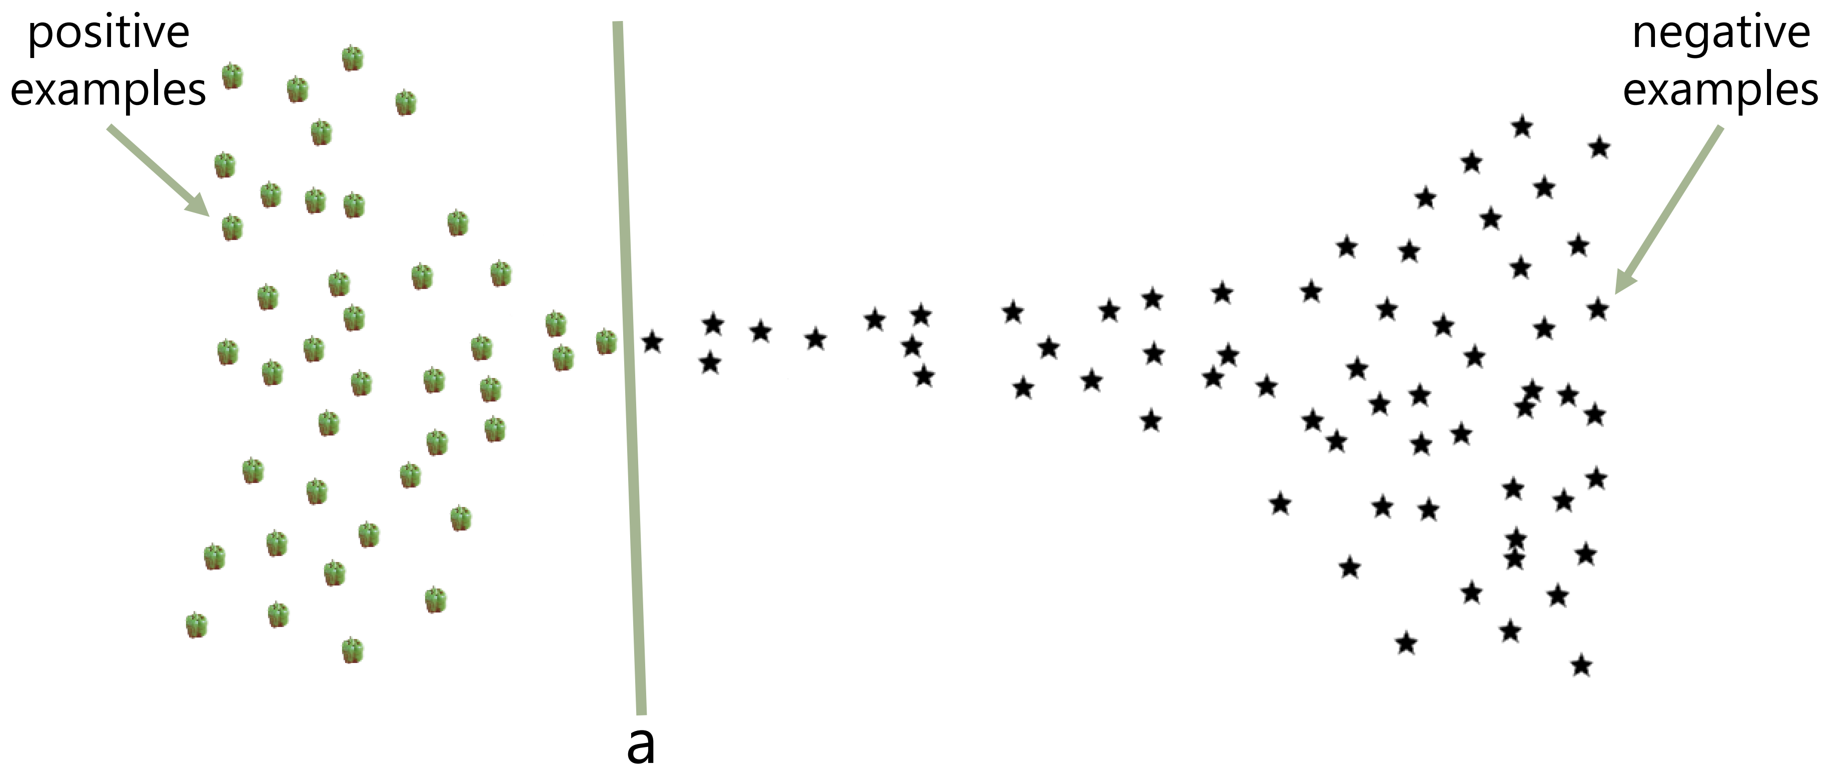

Support Vector Machines

This is essentially the intuition behind Support Vector Machines (SVMs) – train a classifier that focuses on the "difficult" examples by minimizing the misclassification error

We still want a classifier of the form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta - \boxed{\alpha} > 0 \\ \boxed{-1} & \text{otherwise} \end{cases}$$

But we want to minimize the number of misclassifications:

$$\arg\min_\theta \sum_i \delta(y_i(X_i \cdot \theta - \alpha) \leq 0)$$

Support Vector Machines

Simple (seperable) case: there exists a perfect classifier

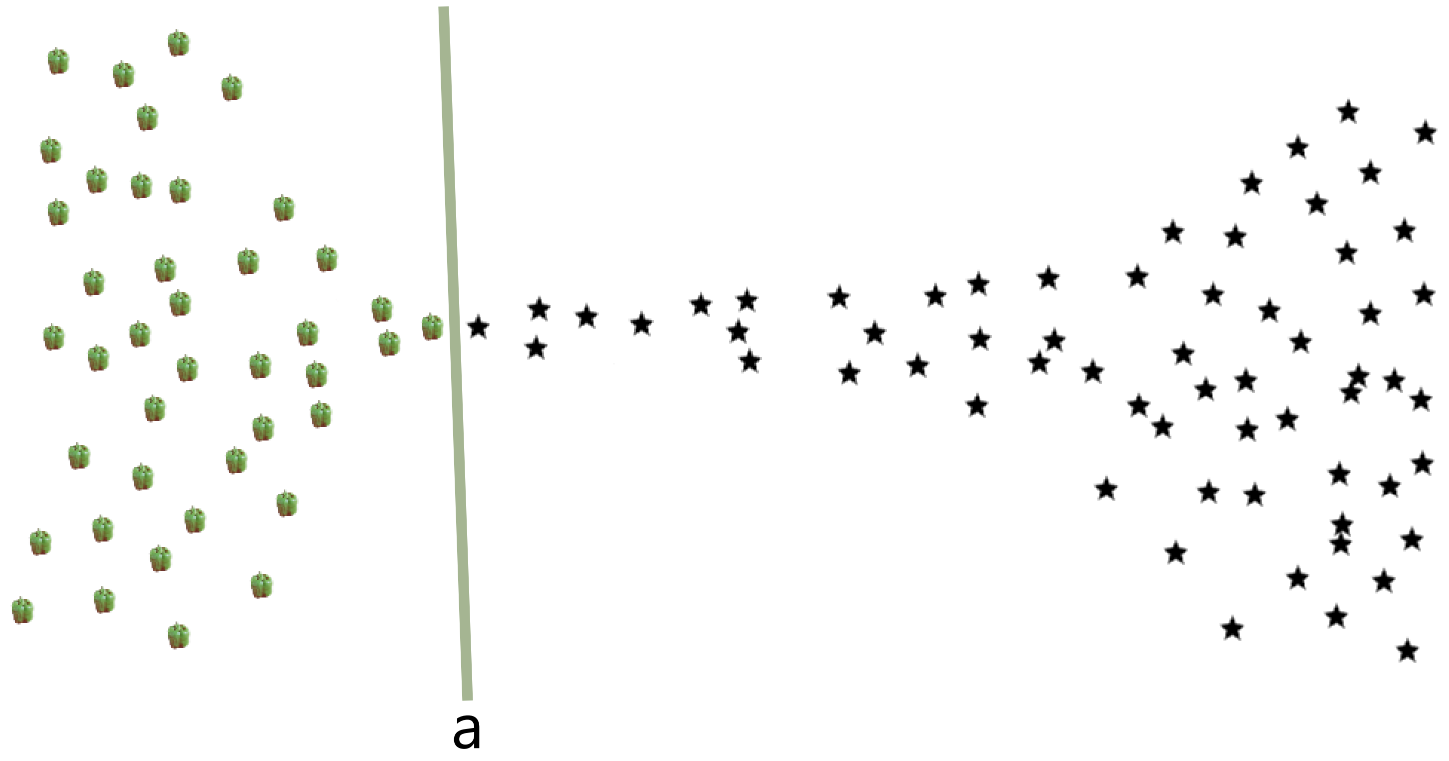

Support Vector Machines

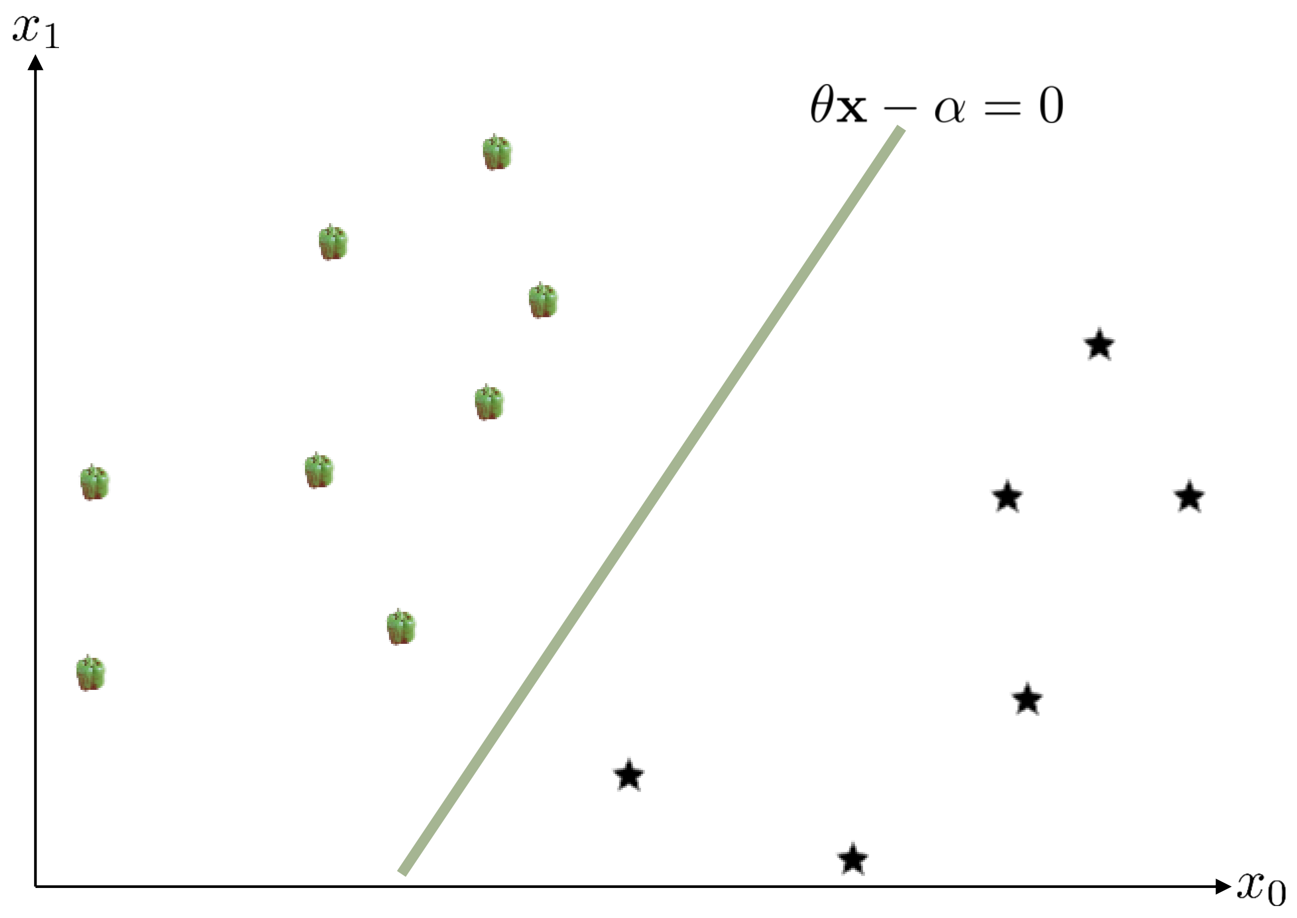

The classifier is defined by the hyperplane $\theta\mathbf{x} - \alpha = 0$

Support Vector Machines

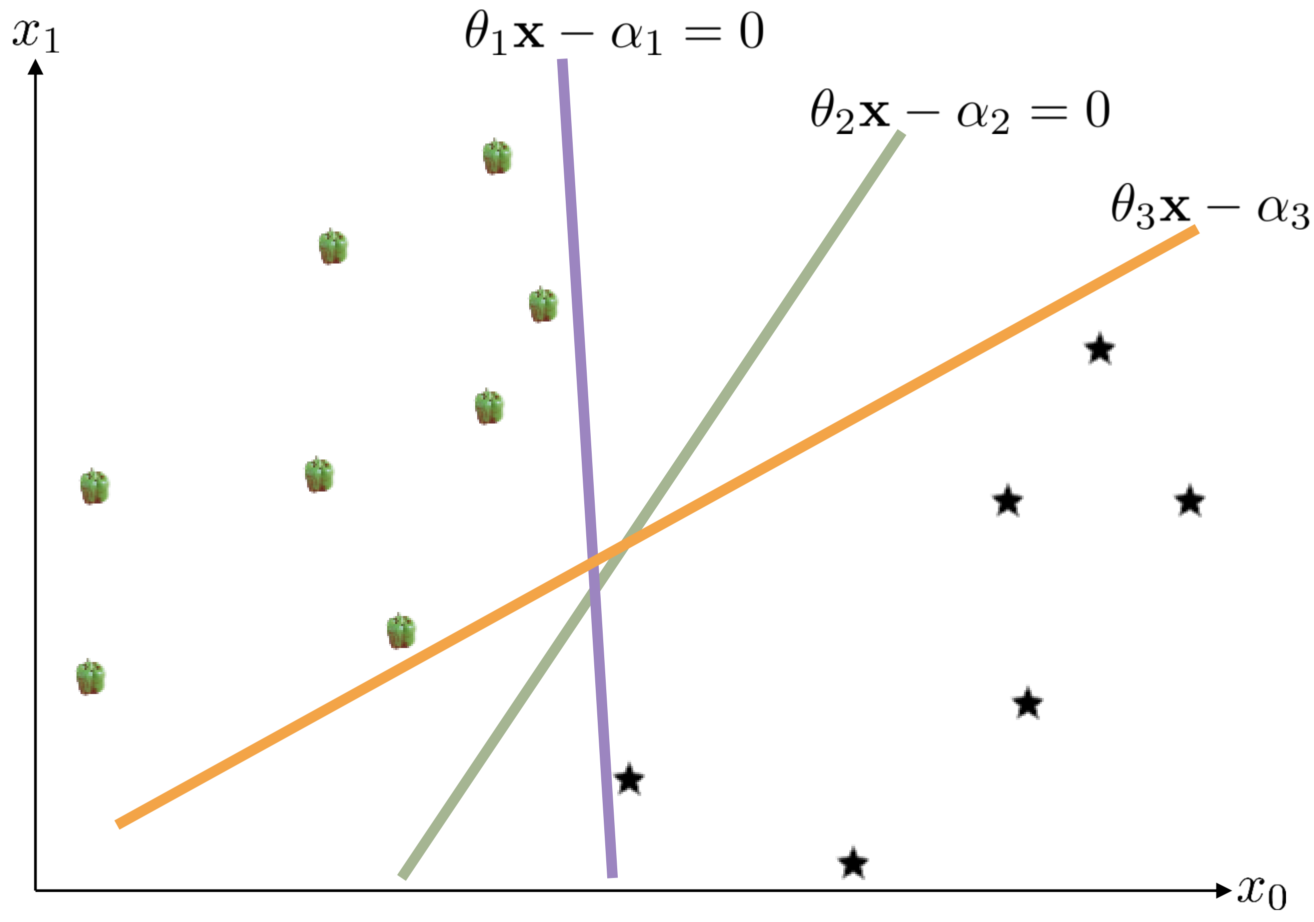
x_1
$\theta_1\mathbf{x} - \alpha_1 = 0$
$\theta_2\mathbf{x} - \alpha_2 = 0$
$\theta_3\mathbf{x} - \alpha_3 = 0$
x_0

Q: Is one of these classifiers preferable over the others?

Support Vector Machines

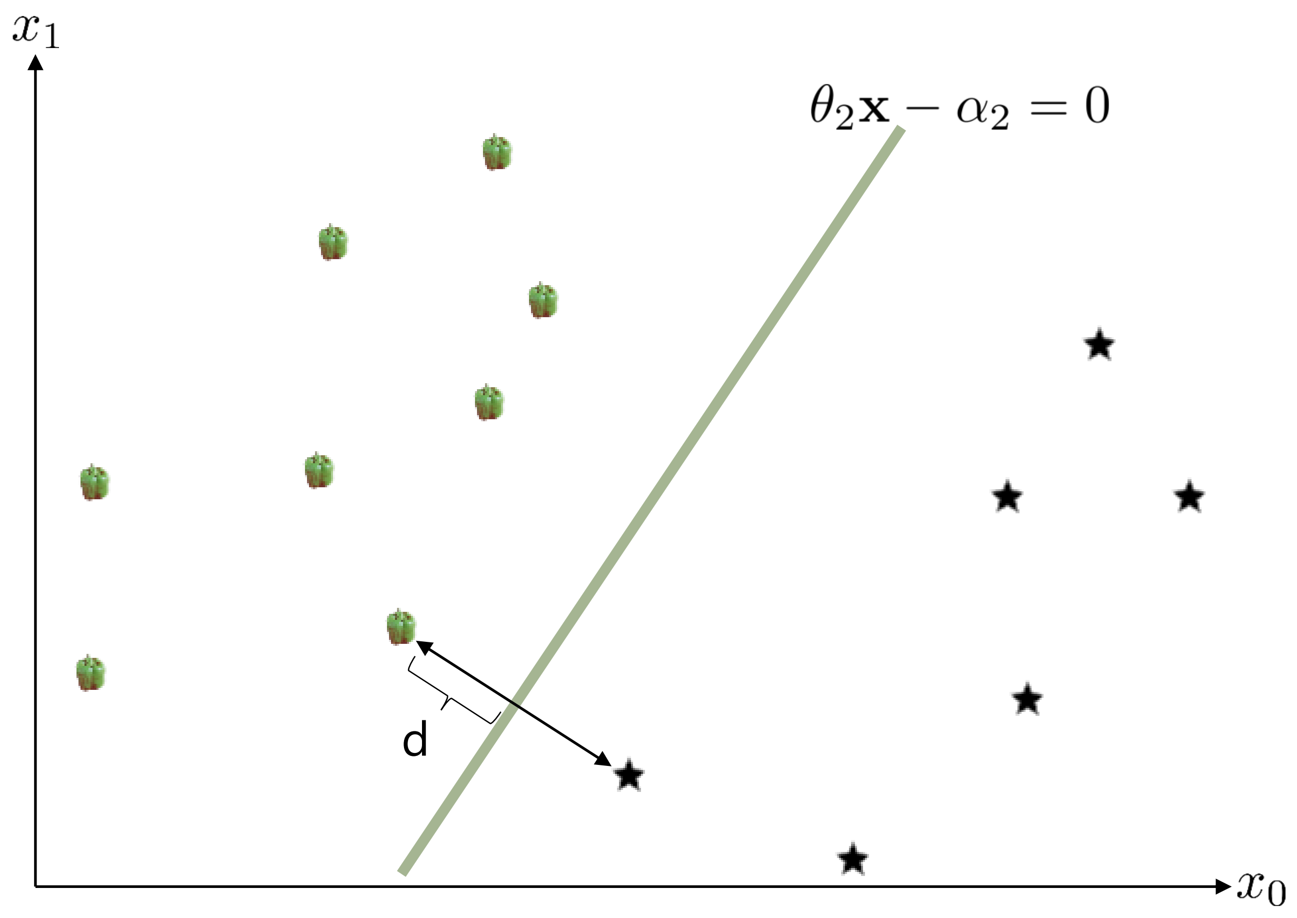

A: Choose the classifier that maximizes the distance to the nearest point

Support Vector Machines

Distance from a point to a line?

Support Vector Machines

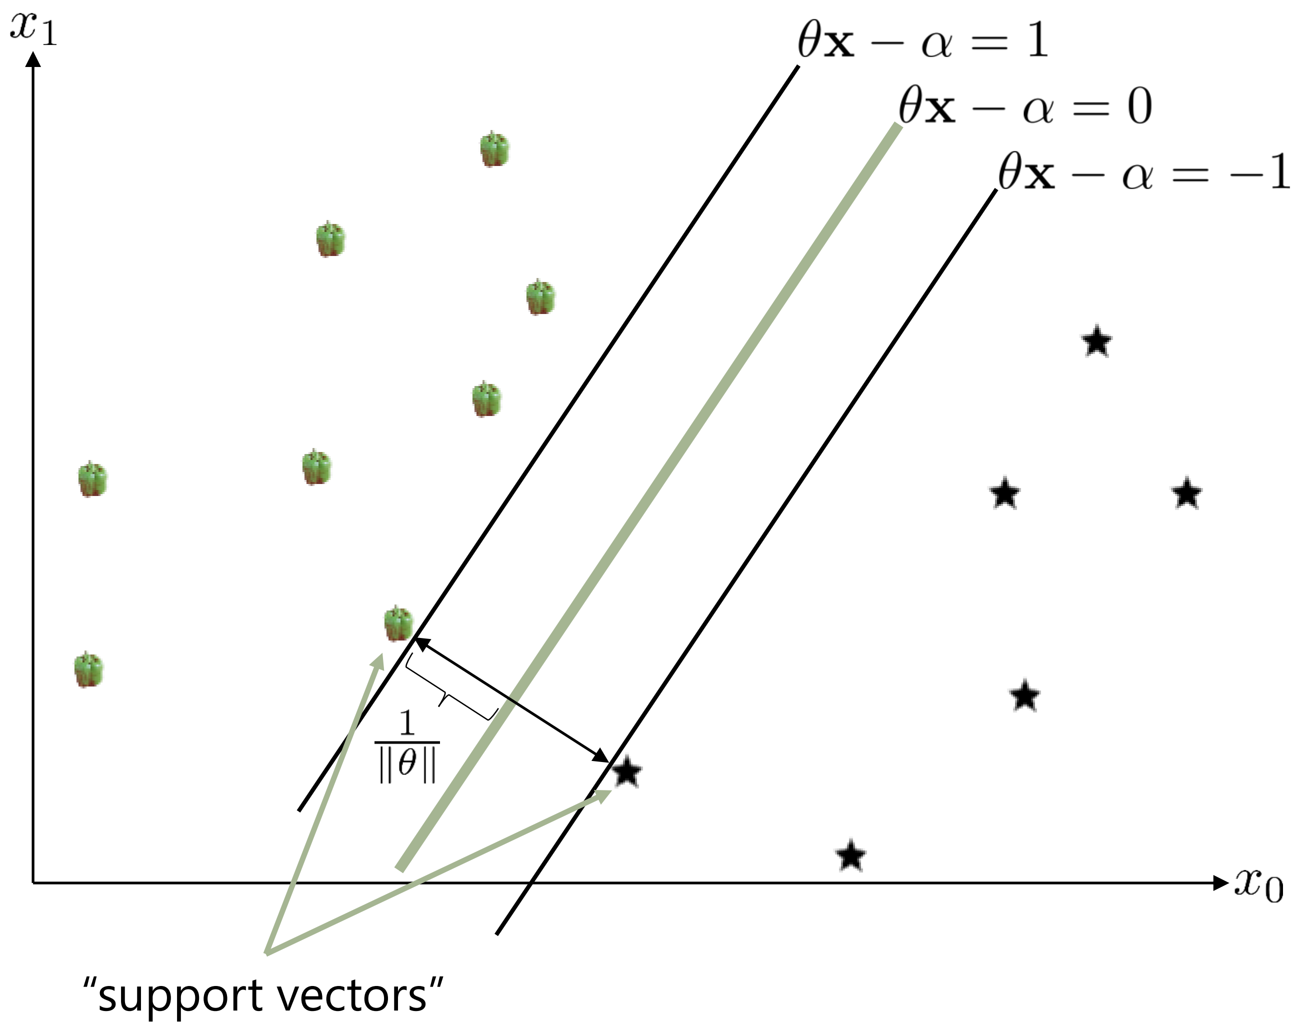
x_1
$\theta\mathbf{x} - \alpha = 1$
$\theta\mathbf{x} - \alpha = 0$
$\theta\mathbf{x} - \alpha = -1$
$\frac{1}{\|\theta\|}$
x_0
"support vectors"

$$\arg\min_{\theta,\alpha} \tfrac{1}{2}\|\theta\|_2^2$$

such that

$$\forall_i y_i(\theta \cdot X_i - \alpha) \geq 1$$

Support Vector Machines

This is known as a "quadratic program" (QP) and can be solved using "standard" techniques

$$\arg\min_{\theta,\alpha} \tfrac{1}{2}\|\theta\|_2^2$$

such that

$$\forall_i y_i(\theta \cdot X_i - \alpha) \geq 1$$

See e.g. Nocedal & Wright ("Numerical Optimization"), 2006

Support Vector Machines

But: is finding such a separating hyperplane even possible?

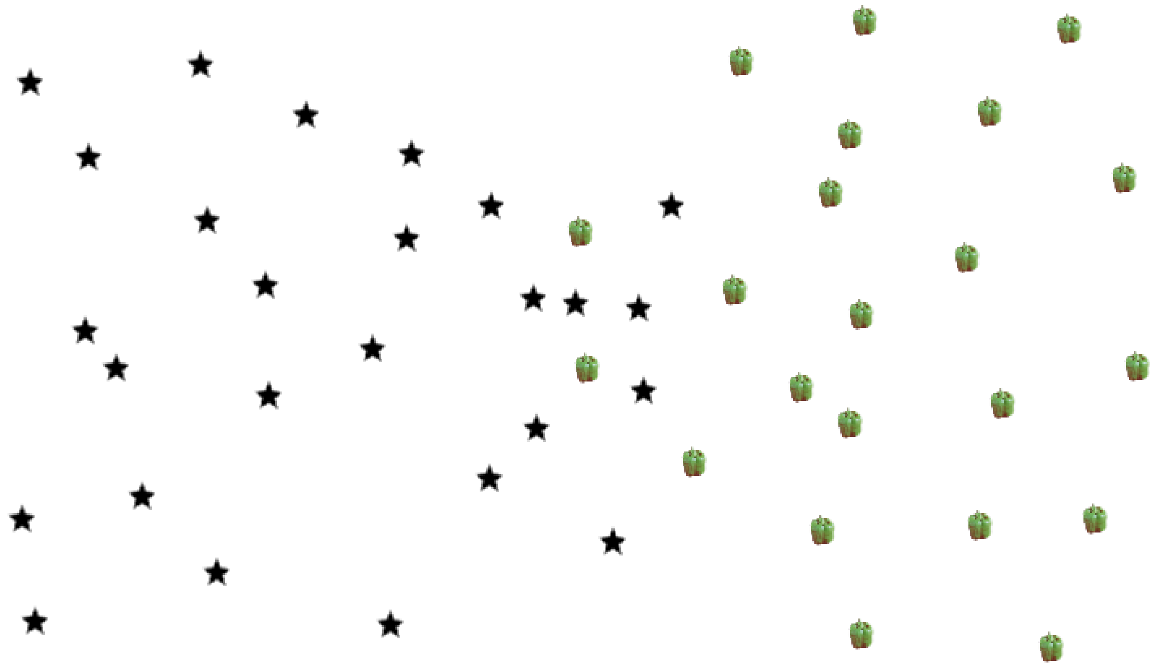

Support Vector Machines

Or: is it actually a good idea?

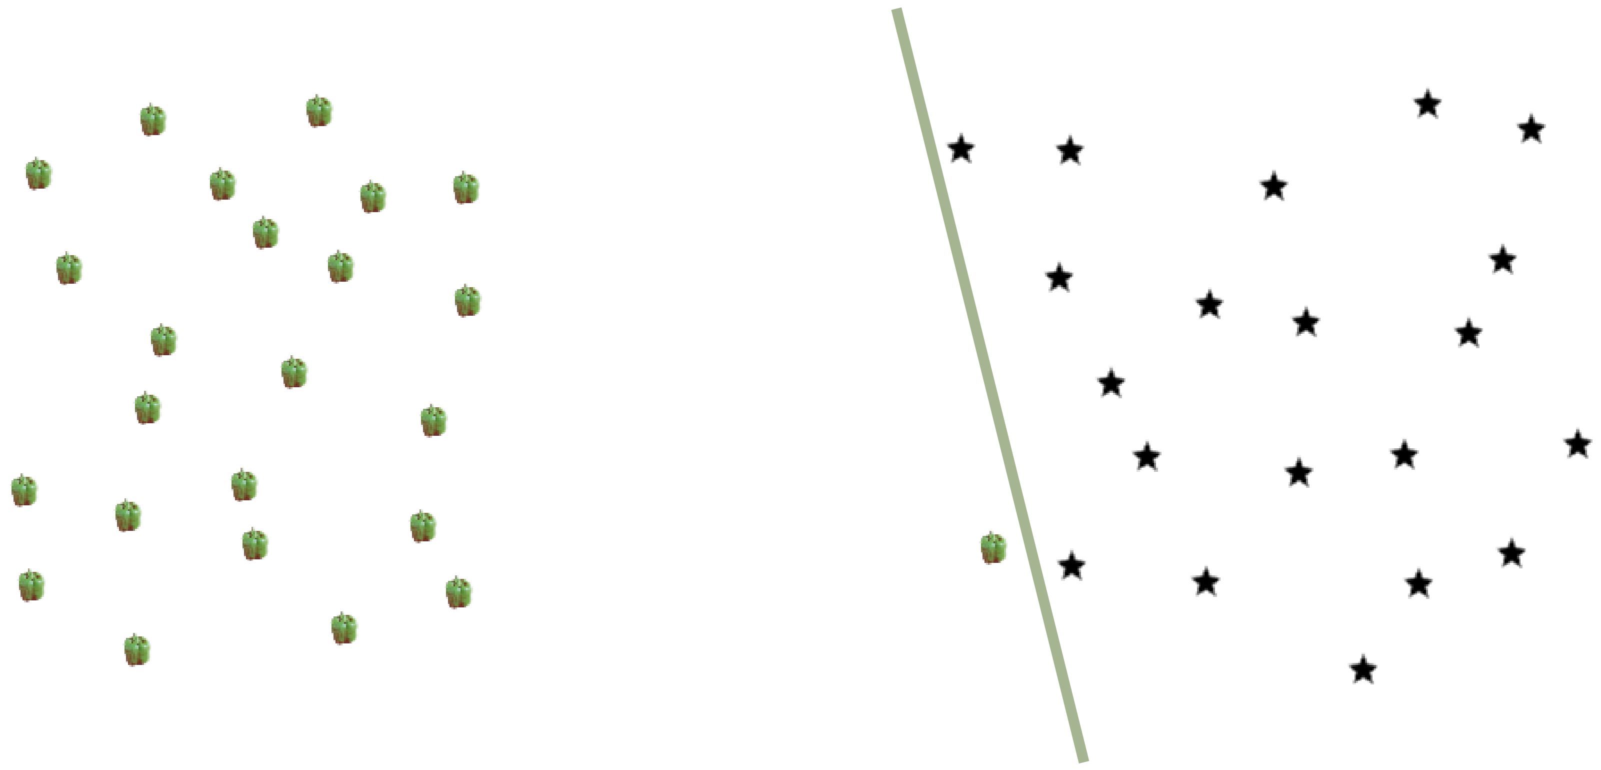

Support Vector Machines

Want the margin to be as wide as possible

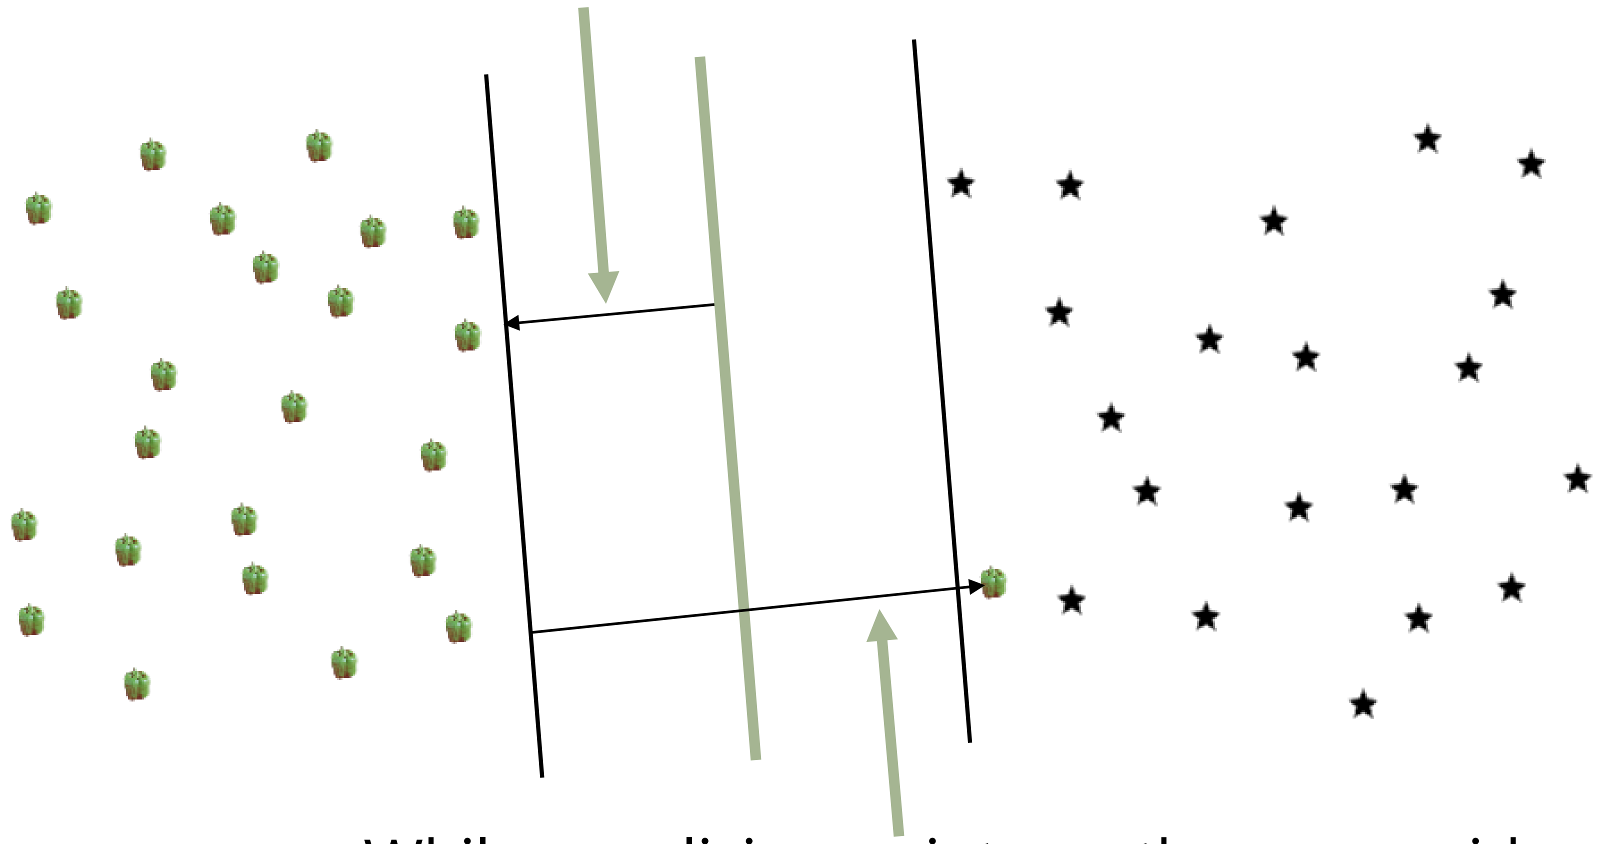

While penalizing points on the wrong side of it

Support Vector Machines

Soft-margin formulation:

$$\arg\min_{\theta,\alpha} \quad \frac{1}{2}||\theta||_2^2$$

such that

$$\forall_i y_i(\theta \cdot X_i - \alpha) \geq 1$$

Summary

The classifiers we've seen this week all attempt to make decisions by associating weights (theta) with features (x) and classifying according to

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Summary

- **Naïve Bayes**
  - Probabilistic model (fits $p(label|data)$)
  - Makes a conditional independence assumption of the form $(feature_i \perp\!\!\!\perp feature_j|label)$ allowing us to define the model by computing $p(feature_i|label)$ for each feature
  - Simple to compute just by counting
- **Logistic Regression**
  - Fixes the "double counting" problem present in naïve Bayes
- **SVMs**
  - Non-probabilistic: optimizes the classification error rather than the likelihood

Pros/cons

- **Naïve Bayes**

  ++ Easiest to implement, most efficient to “train”
  ++ If we have a process that generates feature that *are* independent given the label, it’s a very sensible idea
  -- Otherwise it suffers from a “double-counting” issue

- **Logistic Regression**

  ++ Fixes the “double counting” problem present in naïve Bayes
  -- More expensive to train

- **SVMs**

  ++ Non-probabilistic: optimizes the classification error rather than the likelihood
  -- More expensive to train

Judging a book by its cover

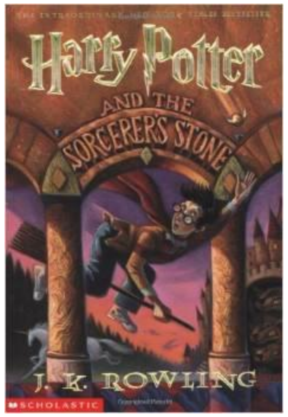

$$\underbrace{[0.723845, 0.153926, 0.757238, 0.983643, \ldots]}_{\text{4096-dimensional image features}}$$

Images features are available for each book on
http://jmcauley.ucsd.edu/cse258/data/amazon/book_images_5000.json

http://caffe.berkeleyvision.org/

Judging a book by its cover

Example: train an SVM to predict whether a book is a children's book from its cover art

(code available on)
http://jmcauley.ucsd.edu/cse258/code/week2.py

Judging a book by its cover

- The number of errors we made was extremely low, yet our classifier doesn’t seem to be very good – why?

CSE 258 – Lecture 4

Web Mining and Recommender Systems

Evaluating classifiers

Which of these classifiers is best?

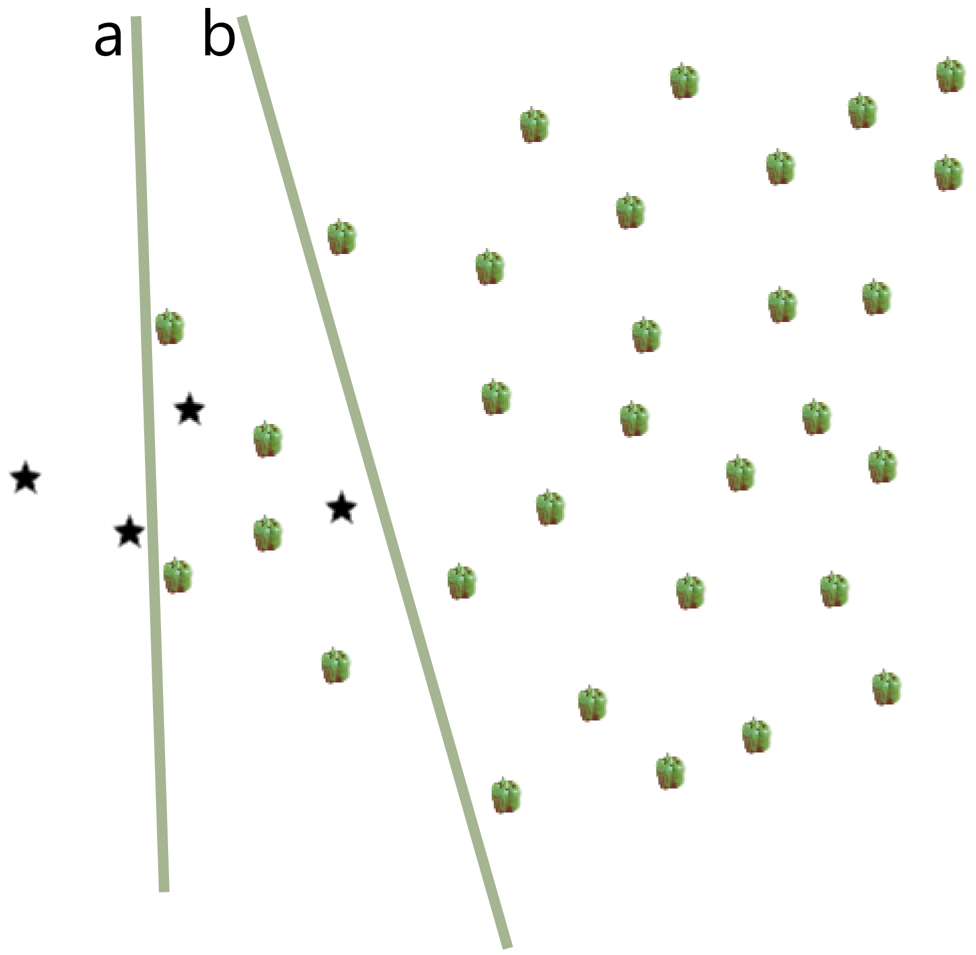

Which of these classifiers is best?

The solution which minimizes the #errors may not be the best one

Which of these classifiers is best?

1. When data are highly imbalanced

If there are far fewer positive examples than negative examples we may want to assign additional weight to negative instances (or vice versa)

e.g. will I purchase a product? If I purchase 0.00001% of products, then a classifier which just predicts "no" everywhere is 99.99999% accurate, but not very useful

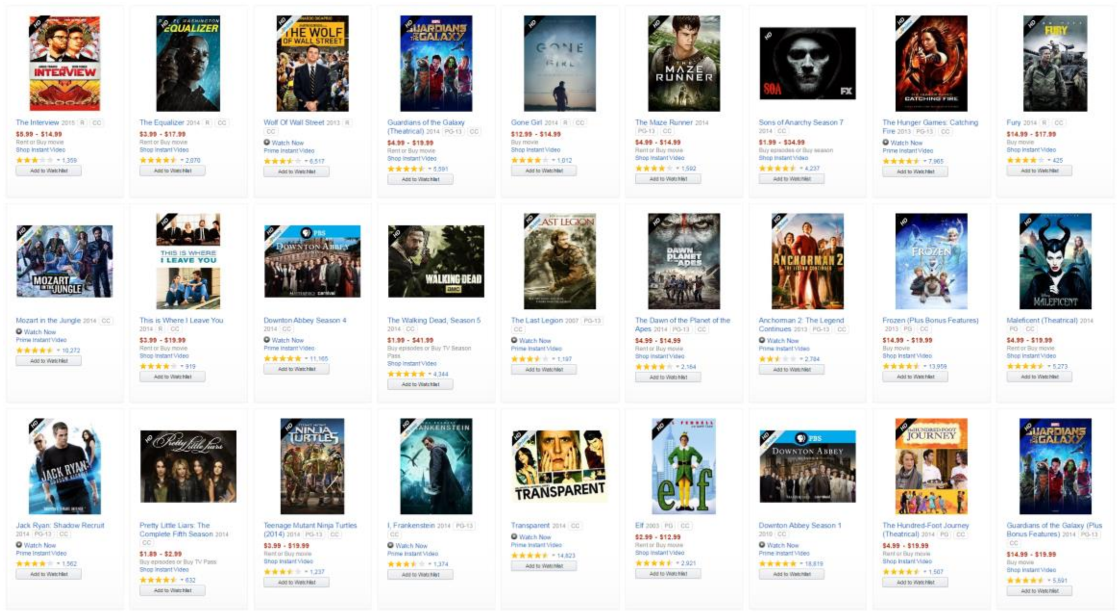

Which of these classifiers is best?

2. When mistakes are more costly in one direction

False positives are nuisances but false negatives are disastrous (or vice versa)

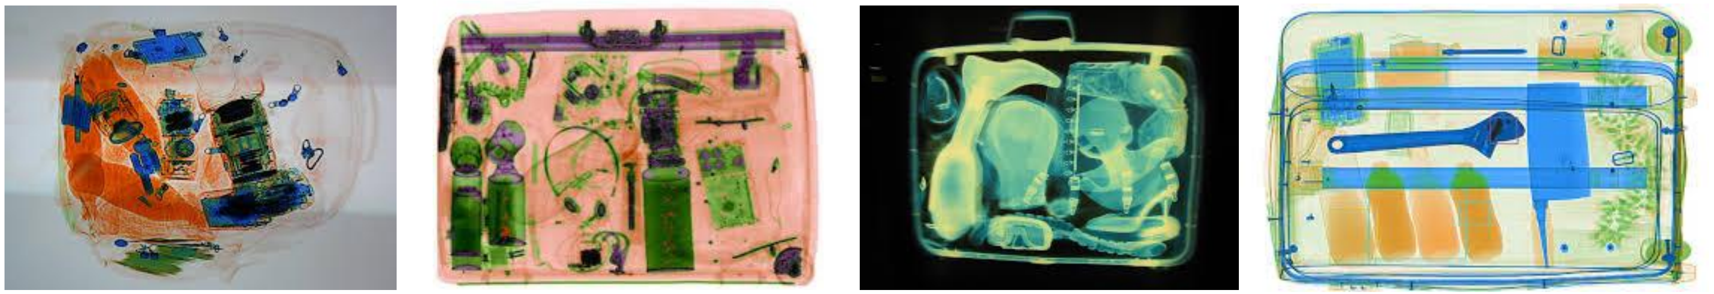

e.g. which of these bags contains a weapon?

Which of these classifiers is best?

3. When we only care about the "most confident" predictions

e.g. does a relevant result appear among the first page of results?

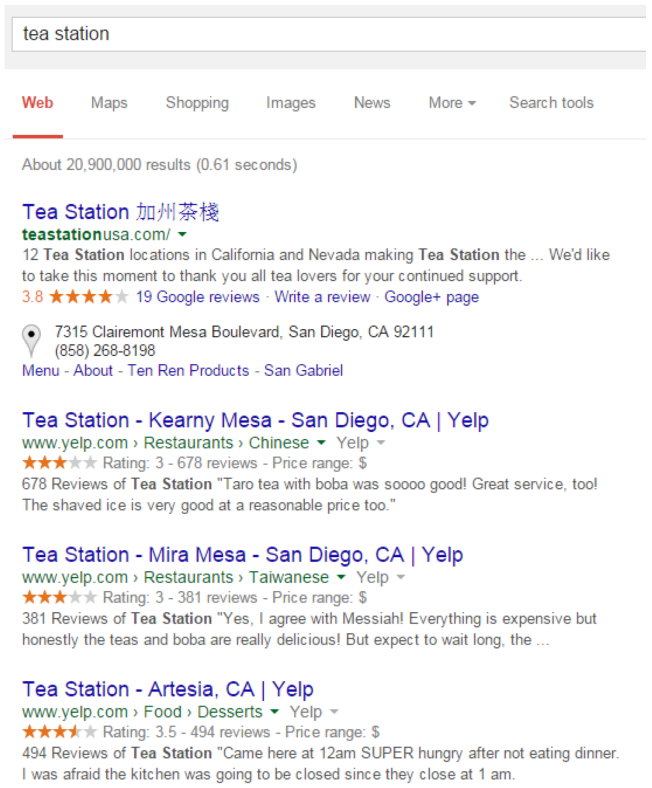

tea station

Web Maps Shopping Images News More Search tools

About 20,900,000 results (0.61 seconds)

Tea Station 加州茶棧
teastationusa.com/
12 **Tea Station** locations in California and Nevada making **Tea Station** the ... We'd like to take this moment to thank you all tea lovers for your continued support.
3.8 ★★★★☆ 19 Google reviews · Write a review · Google+ page

7315 Clairemont Mesa Boulevard, San Diego, CA 92111
(858) 268-8198
Menu - About - Ten Ren Products - San Gabriel

Tea Station - Kearny Mesa - San Diego, CA | Yelp
www.yelp.com › Restaurants › Chinese ▾ Yelp ▾
★★★☆☆ Rating: 3 - 678 reviews - Price range: $
678 Reviews of **Tea Station** "Taro tea with boba was soooo good! Great service, too! The shaved ice is very good at a reasonable price too."

Tea Station - Mira Mesa - San Diego, CA | Yelp
www.yelp.com › Restaurants › Taiwanese ▾ Yelp ▾
★★★☆☆ Rating: 3 - 381 reviews - Price range: $
381 Reviews of **Tea Station** "Yes, I agree with Messiah! Everything is expensive but honestly the teas and boba are really delicious! But expect to wait long, the ...

Tea Station - Artesia, CA | Yelp
www.yelp.com › Food › Desserts ▾ Yelp ▾
★★★½☆ Rating: 3.5 - 494 reviews - Price range: $
494 Reviews of **Tea Station** "Came here at 12am SUPER hungry after not eating dinner. I was afraid the kitchen was going to be closed since they close at 1 am.

Evaluating classifiers

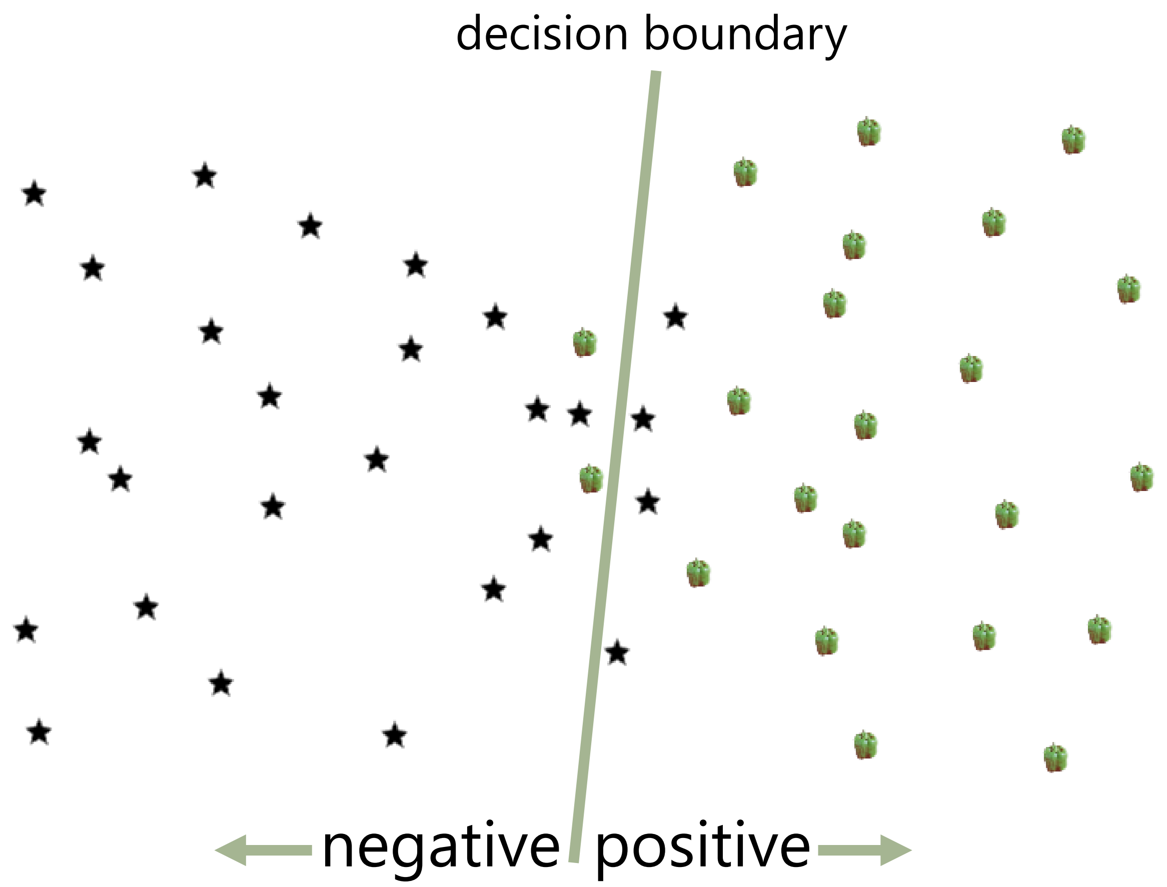

Evaluating classifiers

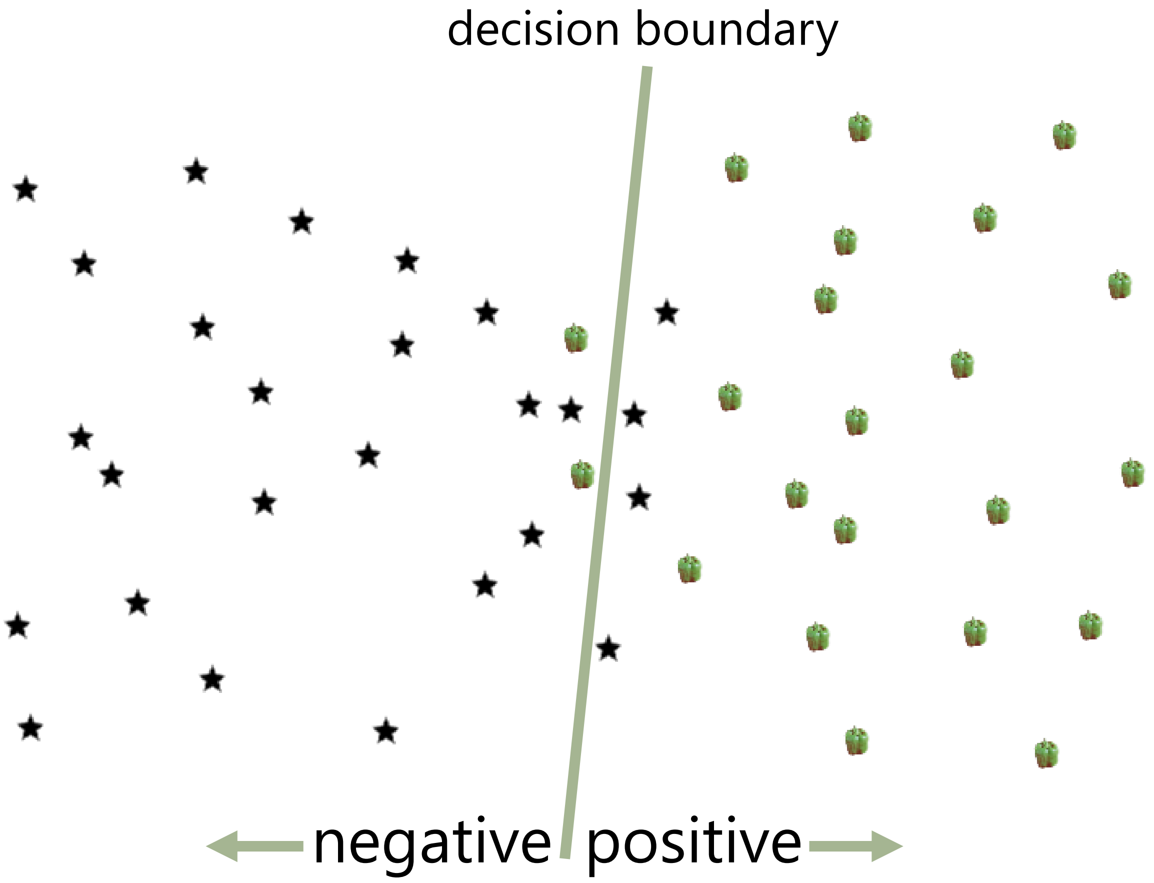

TP (true positive): Labeled as , predicted as

Evaluating classifiers

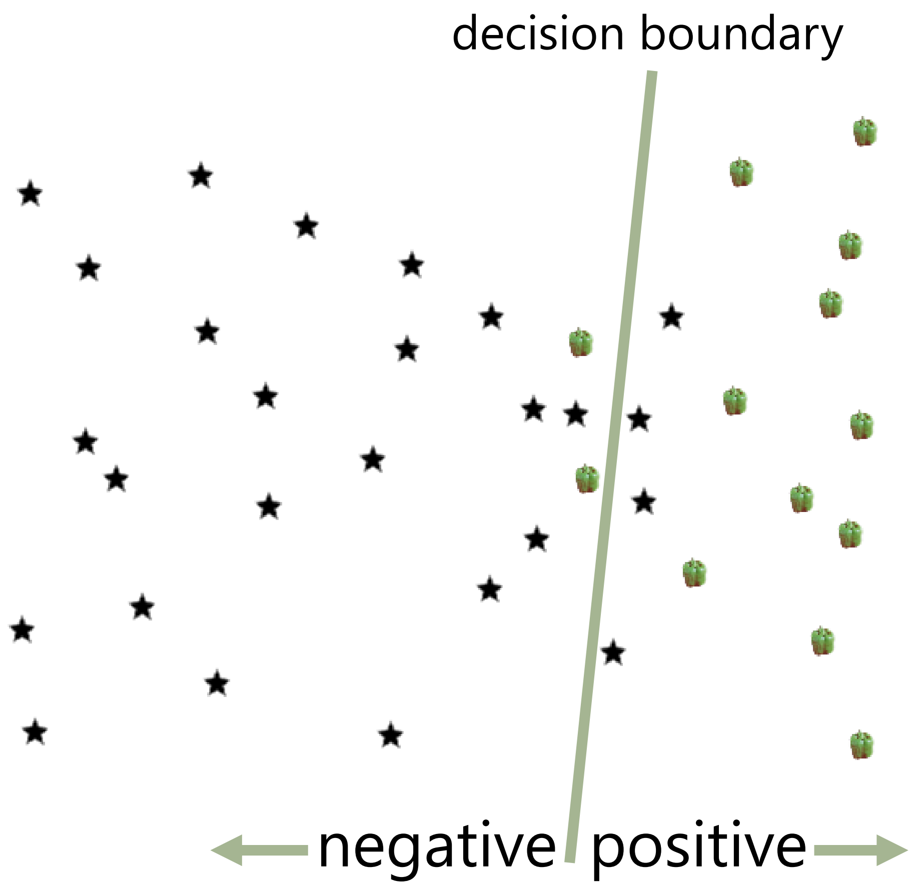

TN (true negative): Labeled as , predicted as

Evaluating classifiers

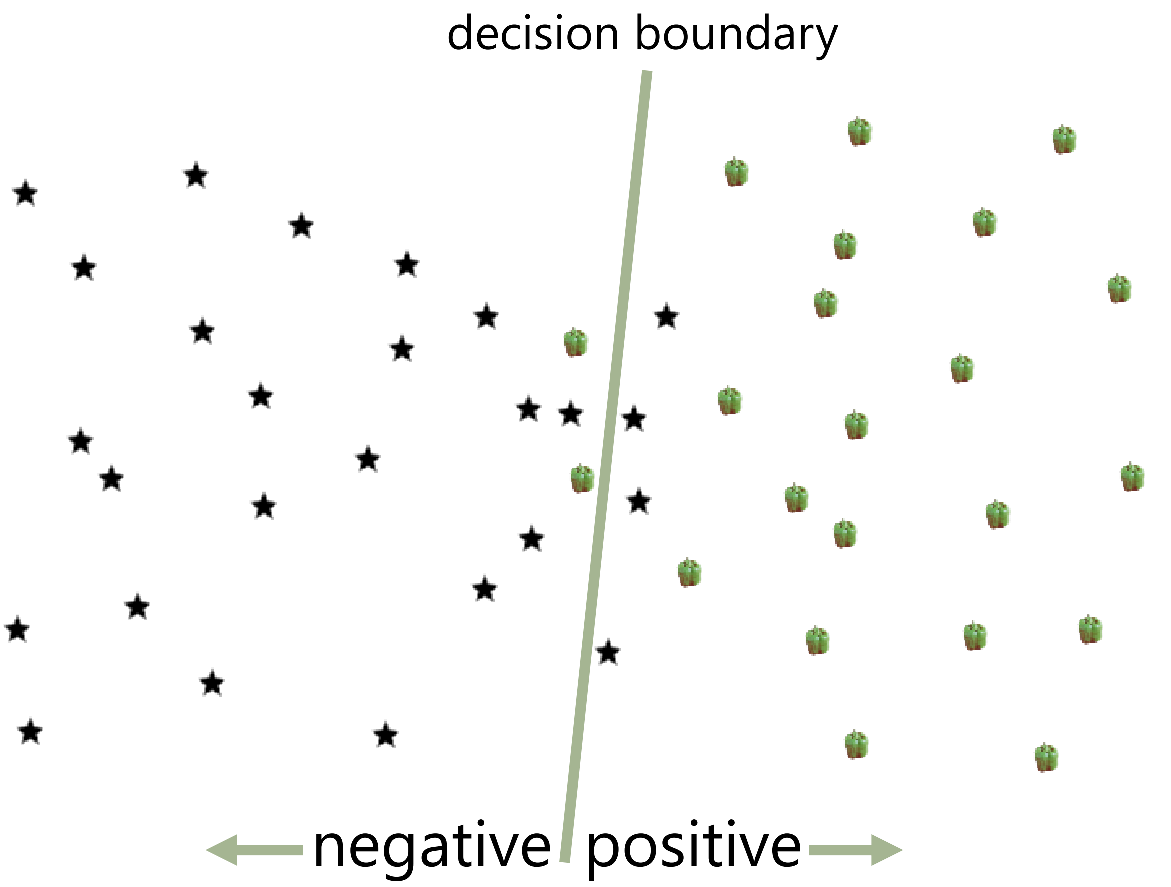

FP (false positive): Labeled as , predicted as

Evaluating classifiers

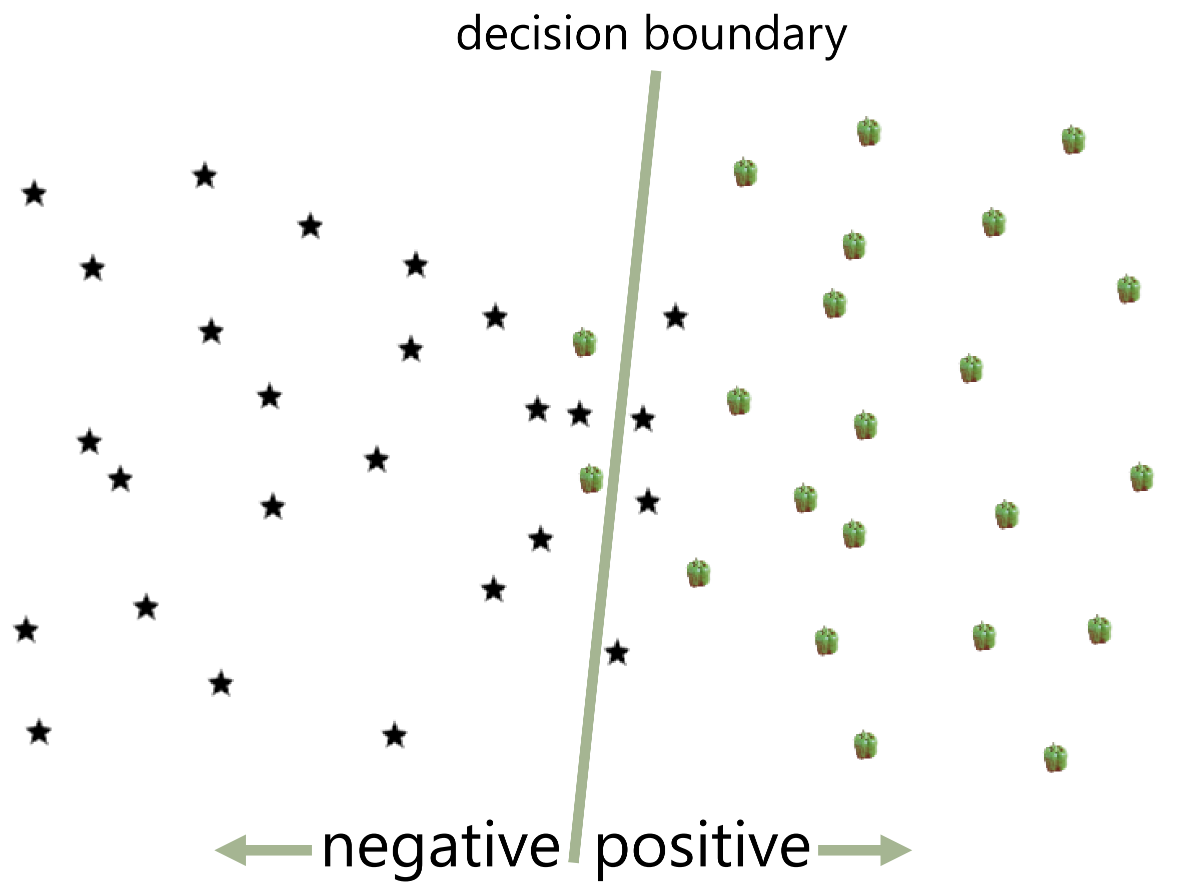

FN (false negative): Labeled as , predicted as

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

Classification accuracy = correct predictions / #predictions

=

Error rate = incorrect predictions / #predictions

=

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

True positive rate (**TPR**) = true positives / #labeled positive
=

True negative rate (**TNR**) = true negatives / #labeled negative
=

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

Balanced Error Rate (BER) = ½ (FPR + FNR)

= ½ for a random/naïve classifier, 0 for a perfect classifier

Evaluating classifiers

e.g.

y = `[  1,  -1,   1,   1,  1, -1,  1,  1,  -1,   1]`

Confidence = `[1.3,-0.2,-0.1,-0.4,1.4,0.1,0.8,0.6,-0.8,1.0]`

Evaluating classifiers

How to optimize a balanced error measure:

$$L_\theta(y|X) = \prod_{y_i=1} p_\theta(y_i|X_i) \prod_{y_i=0}(1 - p_\theta(y_i|X_i))$$

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

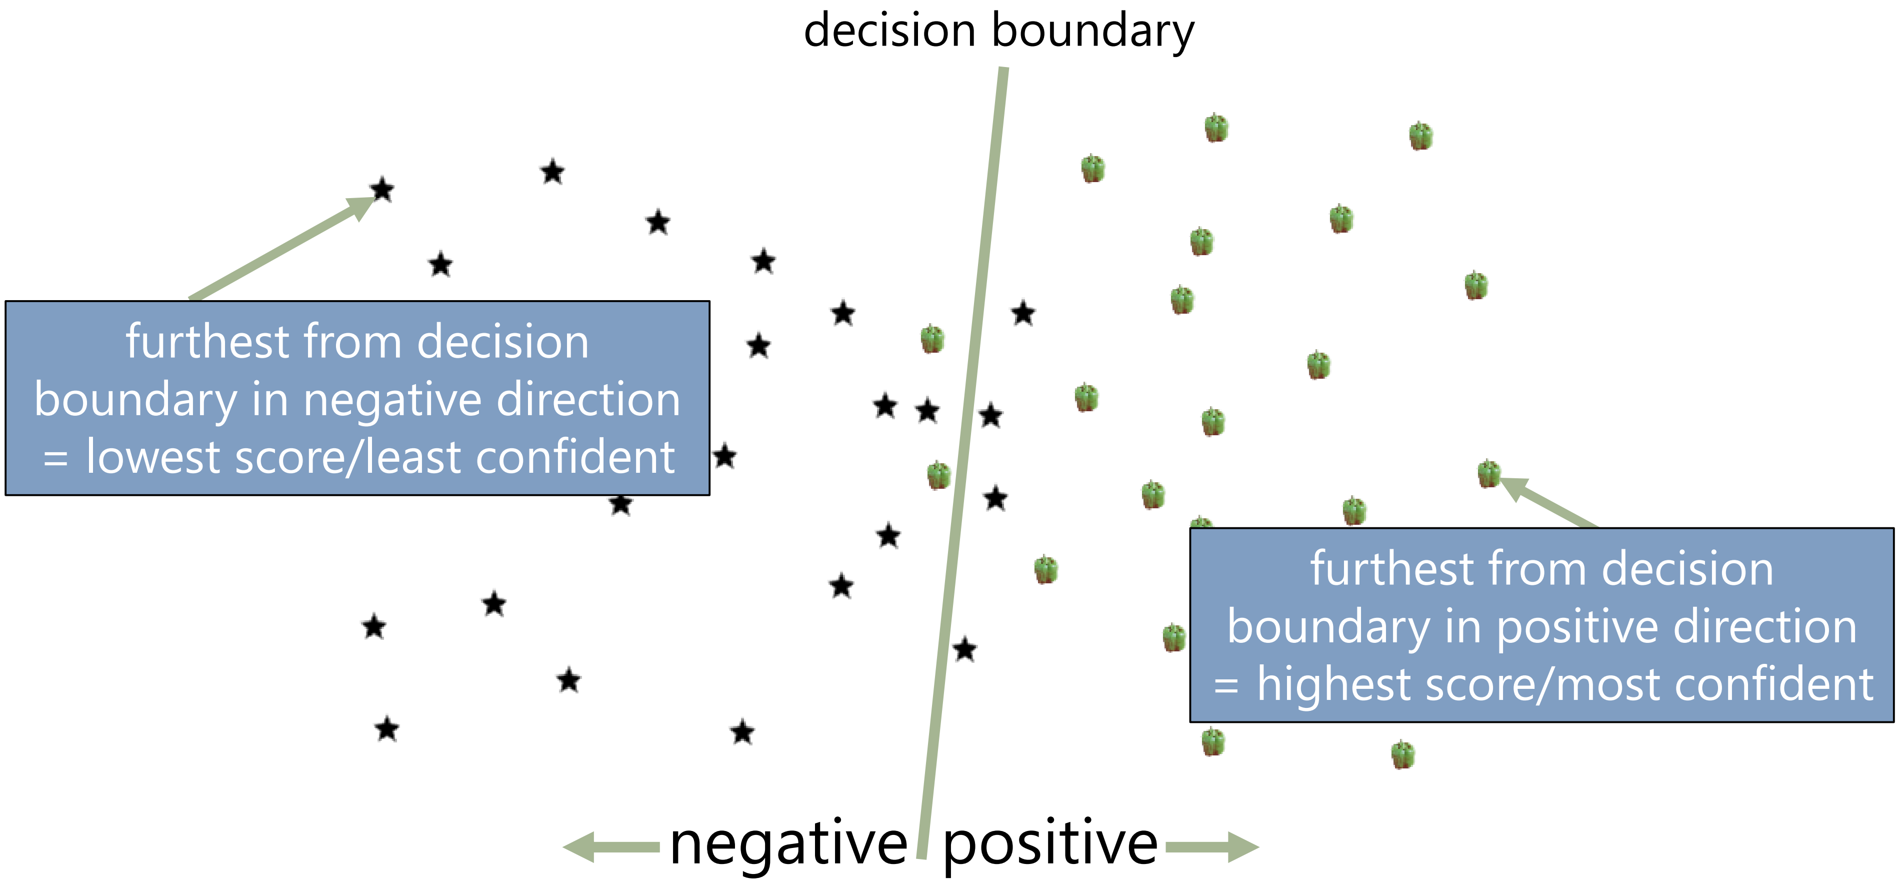
decision boundary
furthest from decision boundary in negative direction = lowest score/least confident
furthest from decision boundary in positive direction = highest score/most confident
negative
positive

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

- In ranking settings, the actual labels assigned to the points (i.e., which side of the decision boundary they lie on) **don't matter**
- All that matters is that positively labeled points tend to be at **higher ranks** than negative ones

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

- For naïve Bayes, the "score" is the ratio between an item having a positive or negative class
- For logistic regression, the "score" is just the probability associated with the label being 1
- For Support Vector Machines, the score is the distance of the item from the decision boundary (together with the sign indicating what side it's on)

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

e.g.

```
         y = [  1,  -1,    1,    1,   1, -1,   1,   1,  -1,   1]
Confidence = [1.3,-0.2,-0.1,-0.4,1.4,0.1,0.8,0.6,-0.8,1.0]
```

Sort **both** according to confidence:

Evaluating classifiers – ranking

The classifiers we’ve seen can associate **scores** with each prediction

Labels sorted by confidence:

$$[\underbrace{1, 1, 1, 1, 1, -1,}\ 1, -1, 1, -1]$$

Suppose we have a fixed budget (say, six) of items that we can return
(e.g. we have space for six results in an interface)

- Total number of **relevant** items =
- Number of items we returned =
- Number of **relevant items** we returned =

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$$\text{precision} = \frac{|\{\text{relevant documents}\} \cap \{\text{retrieved documents}\}|}{|\{\text{retrieved documents}\}|}$$

"fraction of retrieved documents that are relevant"

$$\text{recall} = \frac{|\{\text{relevant documents}\} \cap \{\text{retrieved documents}\}|}{|\{\text{relevant documents}\}|}$$

"fraction of relevant documents that were retrieved"

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$\text{precision@}k$ = precision when we have a budget of *k* retrieved documents

e.g.

- Total number of **relevant** items = 7
- Number of items we returned = 6
- Number of **relevant items** we returned = 5

precision@6 =

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$$F_1 = 2 \cdot \frac{\mathrm{precision} \cdot \mathrm{recall}}{\mathrm{precision} + \mathrm{recall}}$$

(harmonic mean of precision and recall)

$$F_\beta = (1 + \beta^2) \cdot \frac{\mathrm{precision} \cdot \mathrm{recall}}{\beta^2 \mathrm{precision} + \mathrm{recall}}$$

(weighted, in case precision is more important (low beta), or recall is more important (high beta))

Precision/recall curves

How does our classifier behave as we "increase the budget" of the number retrieved items?

- For budgets of size 1 to N, compute the precision and recall
- Plot the precision against the recall

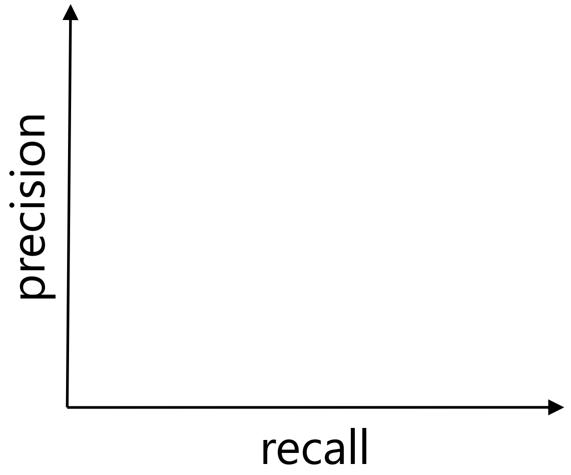

Summary

1. When data are highly imbalanced

If there are far fewer positive examples than negative examples we may want to assign additional weight to negative instances (or vice versa)

e.g. will I purchase product? If I purchase 0.00001% of products, then a classifier which jus predicts "no" everywhere is 99.99999% accurate, but not very useful

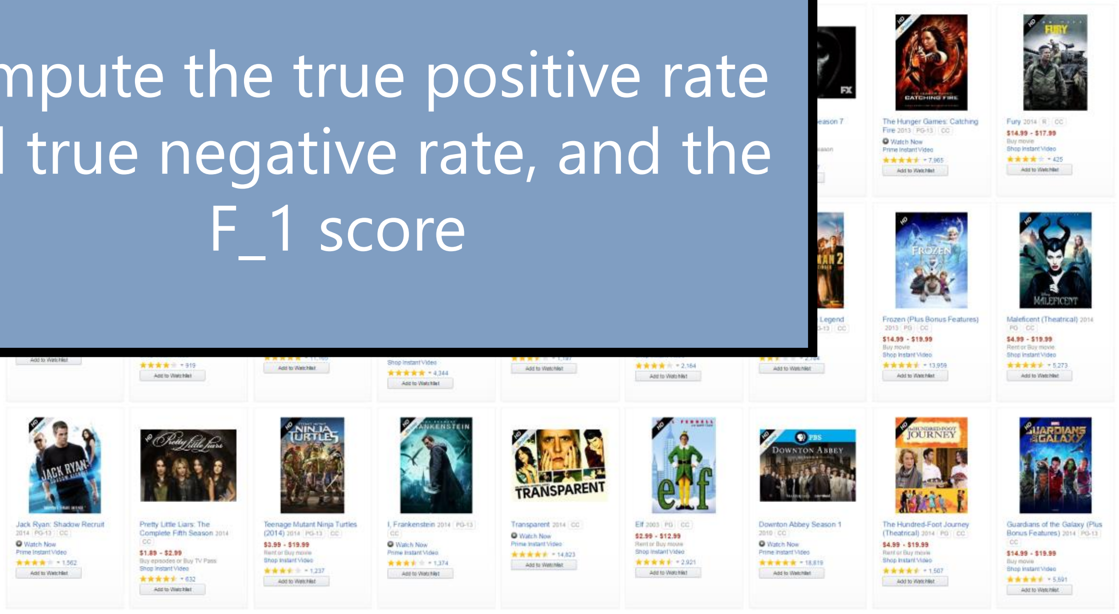

Compute the true positive rate and true negative rate, and the F_1 score

Summary

2. When mistakes are more costly in one direction

False positives are nuisances but false negatives are disastrous (or vice versa)

Compute “weighted” error measures that trade-off the precision and the recall, like the F_\beta score

e.g. which of these bags contains a weapon?

Summary

3. When we only care about the "most confident" predictions

e.g. does
result
among
page of results?

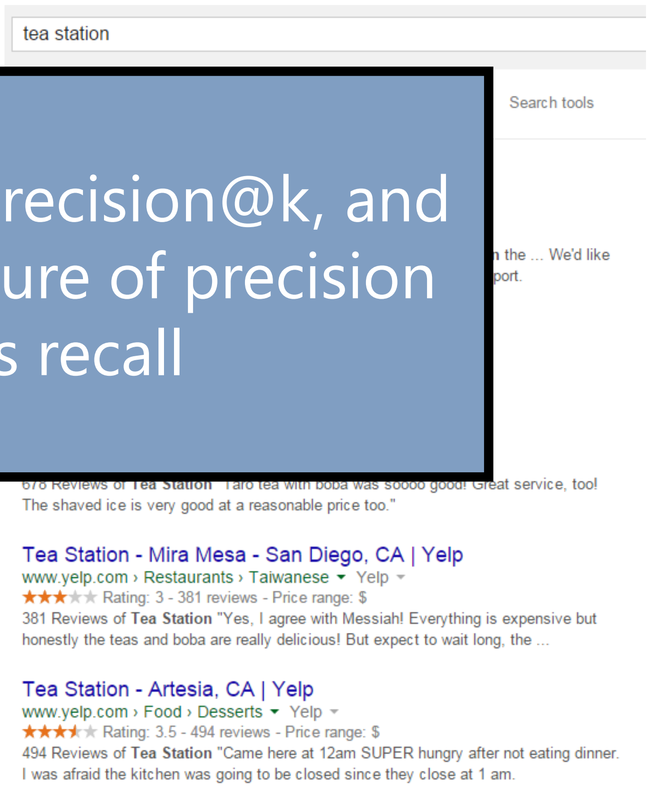

tea station

Search tools

the ... We'd like
port.

678 Reviews of Tea Station "Taro tea with boba was soooo good! Great service, too! The shaved ice is very good at a reasonable price too."

Tea Station - Mira Mesa - San Diego, CA | Yelp
www.yelp.com › Restaurants › Taiwanese ▾ Yelp ▾
★★★☆☆ Rating: 3 - 381 reviews - Price range: $
381 Reviews of Tea Station "Yes, I agree with Messiah! Everything is expensive but honestly the teas and boba are really delicious! But expect to wait long, the ...

Tea Station - Artesia, CA | Yelp
www.yelp.com › Food › Desserts ▾ Yelp ▾
★★★½☆ Rating: 3.5 - 494 reviews - Price range: $
494 Reviews of Tea Station "Came here at 12am SUPER hungry after not eating dinner. I was afraid the kitchen was going to be closed since they close at 1 am.

Compute the precision@k, and plot the signature of precision versus recall

So far: Regression

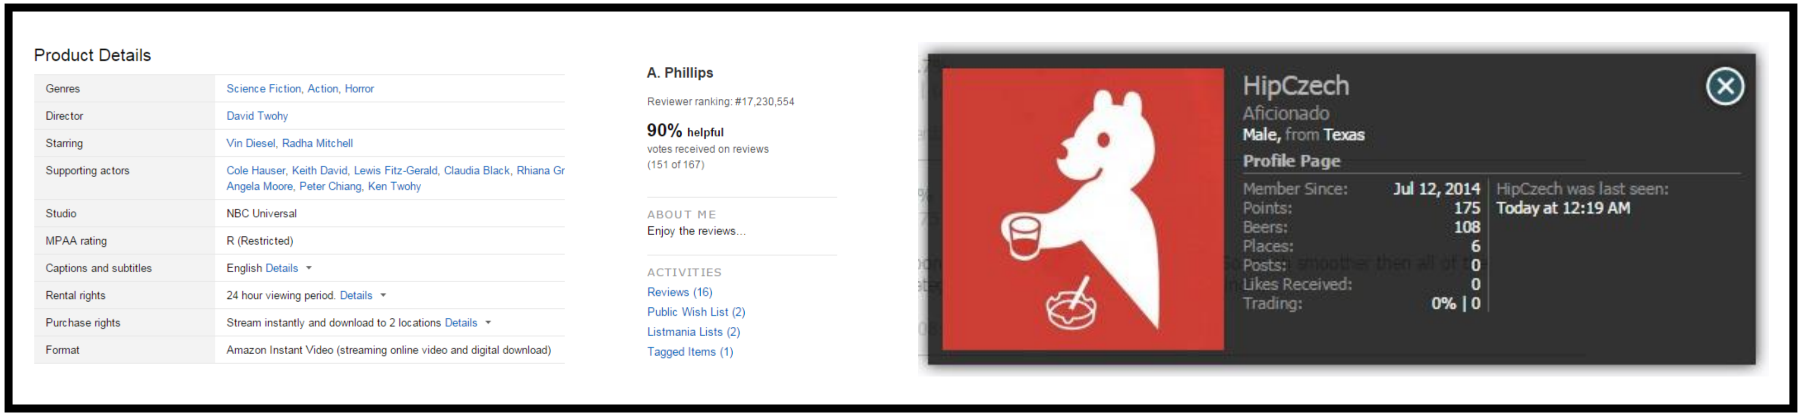

Product Details

Genres	Science Fiction, Action, Horror
Director	David Twohy
Starring	Vin Diesel, Radha Mitchell
Supporting actors	Cole Hauser, Keith David, Lewis Fitz-Gerald, Claudia Black, Rhiana Gr Angela Moore, Peter Chiang, Ken Twohy
Studio	NBC Universal
MPAA rating	R (Restricted)
Captions and subtitles	English Details
Rental rights	24 hour viewing period. Details
Purchase rights	Stream instantly and download to 2 locations Details
Format	Amazon Instant Video (streaming online video and digital download)

A. Phillips

Reviewer ranking: #17,230,554

90% helpful

votes received on reviews (151 of 167)

ABOUT ME

Enjoy the reviews...

ACTIVITIES

Reviews (16)

Public Wish List (2)

Listmania Lists (2)

Tagged Items (1)

HipCzech

Aficionado

Male, from Texas

Profile Page

Member Since:	Jul 12, 2014	HipCzech was last seen:
Points:	175	Today at 12:19 AM
Beers:	108	
Places:	6	
Posts:	0	
Likes Received:	0	
Trading:	0% | 0	

How can we use **features** such as product properties and user demographics to make predictions about **real-valued** outcomes (e.g. star ratings)?

How can we prevent our models from **overfitting** by favouring simpler models over more complex ones?

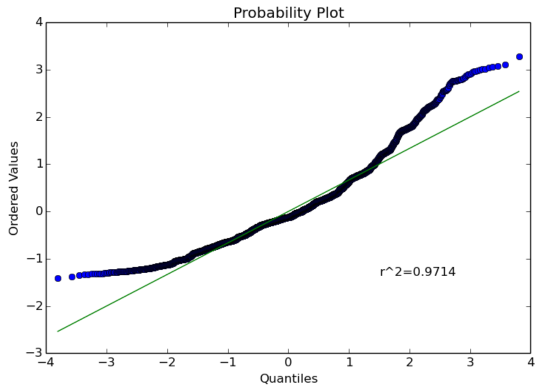

How can we assess our decision to optimize a particular error measure, like the MSE?

So far: Classification

Next we adapted these ideas to **binary** or **multiclass** outputs

What animal is in this image?

Pitch Black - Unrated Director's Cut R

777 IMDb 7.1/10

Watch Trailer

When their ship crash-lands on a remote pla
escaped convict Riddick (Vin Diesel) isn't th
lurk in the shadows, waiting to attack in the

See More

Starring: Vin Diesel, Radha Mitchell
Runtime: 1 hour, 53 minutes
Available to watch on supported devices.

Will I **purchase** this product?

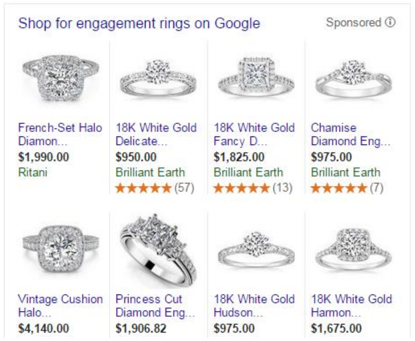
Shop for engagement rings on Google Sponsored

French-Set Halo Diamon... $1,990.00 Ritani	18K White Gold Delicate... $950.00 Brilliant Earth ★★★★★ (57)	18K White Gold Fancy D... $1,825.00 Brilliant Earth ★★★★★ (13)	Chamise Diamond Eng... $975.00 Brilliant Earth ★★★★★ (7)
Vintage Cushion Halo... $4,140.00	Princess Cut Diamond Eng... $1,906.82	18K White Gold Hudson... $975.00	18K White Gold Harmon... $1,675.00

Will I **click on** this ad?

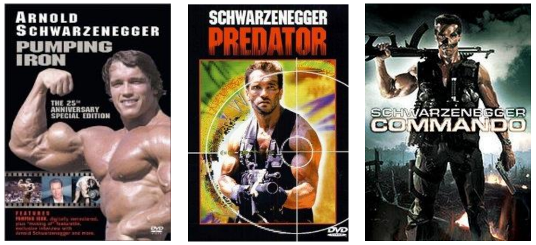

Combining features using naïve Bayes models

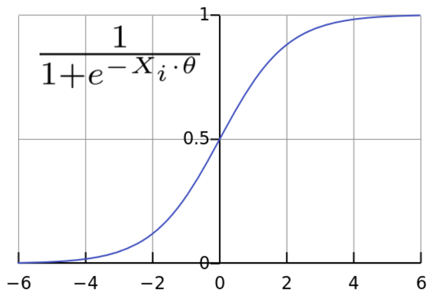

Logistic regression

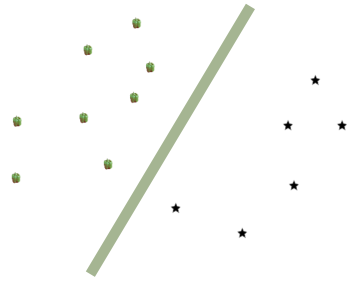

Support vector machines

So far: supervised learning

Given **labeled training data** of the form

$$\{(\mathrm{data}_1, \mathrm{label}_1), \ldots, (\mathrm{data}_n, \mathrm{label}_n)\}$$

Infer the function

$$f(\mathrm{data}) \overset{?}{\rightarrow} \mathrm{labels}$$

So far: supervised learning

We've looked at two types of prediction algorithms:

Regression $\rightarrow$

$$y_i = X_i \cdot \theta$$

Classification $\searrow$

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Questions?

Further reading:

- "Cheat sheet" of performance evaluation measures: http://www.damienfrancois.be/blog/files/modelperfcheatsheet.pdf
- Andrew Zisserman's SVM slides, focused on computer vision: http://www.robots.ox.ac.uk/~az/lectures/ml/lect2.pdf